

Relational Algebra

Definitions

- ↳ attribute \Rightarrow cols. in a relation
- ↳ tuples \Rightarrow rows in a relation
- ↳ domain \Rightarrow type of an attr.
- ↳ schema \Rightarrow the structure of a relation
- ↳ instance \Rightarrow content of a relation
- ↳ key \Rightarrow attr. that uniquely identify a tuple

Uses SET SEMANTICS

- ↳ No dups + no ordering

Operators

↳ σ_R

↳ $\pi_{AR} \Rightarrow \pi_A(\pi_R)$

↳ $R \times S$

- ↳ output per pair of input tuples

↳ SET SEMANTICS

↳ $R \bowtie S$

- ↳ Not comm., generalizes into \times

↳ $R \bowtie R = R$

↳ $\rho_{SR} \Rightarrow$ rename

↳ $R \cup S, R - S, R \cap S$

- ↳ Should have same schema

↳ \cap not comm. $\Rightarrow R \cap S = R - (R - S)$

Basic SQL

Types

↳ CHAR(N)

↳ VARCHAR(N)

↳ INT

↳ DECIMAL(N,M)

- ↳ N sig figs, M after dec.

↳ REAL

↳ DOUBLE

↳ TIMESTAMP

TABLE creation

↳ CREATE TABLE <name domain, ...>

- ↳ PRIMARY KEY (attr) \Rightarrow can't be NULL

↳ UNIQUE (attr)

↳ DEFAULT <val>

↳ DROP TABLE <schema>

Data manipulation

↳ INSERT INTO <schema>

VALUES (<val> | <relation>)

↳ DELETE FROM <schema>

WHERE <cond>

↳ UPDATE <schema>

SET A = A, ...

WHERE <cond>

Operators

↳ SELECT

↳ + DISTINCT

↳ FROM

↳ WHERE

↳ UNION, INTERSECT, EXCEPT

↳ SET SEMANTICS

↳ + ALL for bag semantics

↳ NOT IN, IN

↳ ALL, SOME

↳ Ex: $a \neq \text{NULL}$

Subqueries

↳ Nested SELECT

↳ $1 \times 1 \Rightarrow$ scalar

Correlated subquery

↳ References a table in the outer clause

↳ EXISTS ()

Common Table Expr.

↳ WITH alias > AS subquery

...

ORDER BY/FETCH FIRST

↳ ORDER BY <attr> ASC/DESC

↳ OFFSET > FETCH FIRST <num> ROWS ONLY

↳ changes bag semantics

↳ LIMIT <num> OFFSET <offset>

General SELECT

↳ SELECT ...

FROM ...

WHERE ...

GROUP BY ...

HAVING ...

ORDER BY ...

FETCH FIRST ...

NULLS

↳ NULLS arithmetic ops = NULL

↳ $0 \text{ DIV } 0 = \text{NULL}$

↳ NULL & NULL

↳ logical ops \neq NULL = Unknown

↳ Aggregates ignore tuples \neq NULL

↳ NOT COUNT(*)

↳ AGG(NULL) = NULL

↳ COUNT(*) = 0

↳ Set ops treat NULL specially

↳ IS NULL or IS NOT NULL

↳ COALESCE() \Rightarrow 1st non NULL

Outer join

↳ preserves dangling tuples

↳ LEFT or RIGHT

Recursion

↳ WITH RECURSIVE

procedure (data, ancestor)

AS (SELECT ...

FROM Parent)

(SELECT P.data, A.ancestor

FROM Parent P, ancestor A

WHERE operator P.child)

SELECT ancestor

FROM Ancestor

WHERE data = 'Succ.';

Advanced SQL

Aggregate func.

↳ AVG, SUM, COUNT, MIN, MAX

↳ results \Rightarrow single tuple

↳ shouldn't appear in WHERE

↳ GROUP BY: forms groups where a single tuple exists for each unique attr.

↳ SELECT has Agg. Function on unique attr. only

↳ HAVING: checks aggregate conds.

Window Func.

↳ FUNCTION (attr) OVER()

↳ Applies agg. over all inputs

↳ PARTITION BY: agg. over partition

Case Expr.

↳ CASE <attr> WHEN <cond> THEN <expr> ELSE <expr> END

E/R Model

- ↳ 2 sets: entity and relationship p
 - ↳ entity:
 - ↳ relationship:
- ↳ Cardinality
 - ↳ # entities participate
 - ↳ 1-1, 1-many, many-many
 - ↳ minimum "1" side
 - ↳ total participation =
 - ↳ General notation
 - ↳ a..b, * = infinite
- ↳ Roles: designated to edges from an entity set
 - ↳ label the role of that entity
- ↳ Superclass and Subclass
 - ↳ Specialization and Generalization
 - ↳ Subclass inherits all attr of superclass
 - ↳ Subclass participates in the relationship of the superclass
 - ↳ Subclass may participate in its own relationship
 - ↳ Disjoint specialization vs. overlapping specialization
 - ↳ Either-or vs. multiple specialization
- ↳ or
- ↳ Weak Entity set
 - ↳ not enough attributes
 - ↳
 - ↳ takes comes from linked entity sets
 - ↳ identifying relationship: weak entity set and owner
 - ↳

MapReduce

- Model
 - ↳ Programmer primitives
 - ↳ Map function: $\text{unit_data} \rightarrow (k, v), (k', v'), \dots$
 - ↳ Reduce function: $(k, v_1), (k, v_2), \dots \rightarrow (k, \text{agg}(v_1, v_2, \dots))$
 - ↳ MapReduce handles the rest
 - ↳ Automatic data partition, distribution, and collection
 - ↳ Failure and speed-disparity handling
- Spark
 - ↳ map(): convert 1 input tuple into output tuple
 - ↳ flatMap(): convert 1 input tuple into multiple outputs
 - ↳ reduceByKey(): specifies how 2 input values agg.
 - ↳ filter(): filter out tuples based on cond.

Normalization

FD

- ↳ $U[X] \Rightarrow U[\text{sid}, \text{name}] \Rightarrow (\text{sid}, \text{name})$
- ↳ $X \rightarrow Y \Rightarrow$ no tuple can have $X_1 \neq X_2$ but $Y_1 \neq Y_2$
- ↳ Trivial: $(\text{sid}) \Rightarrow \text{sid}$
- Logical implication
 - ↳ $A \rightarrow B, B \rightarrow C \Rightarrow A \rightarrow C$
- Closure
 - ↳ F^+ is the set of FDs logically implied by F
 - ↳ F^+ is set of all attributes that are functionally determined by X
 - ↳ start with $X \rightarrow X$ and repeat
 - ↳ X is a key if $X \rightarrow$ all attr
- Lossless Decomp. is when the shared attr. are a key of one of the tables
- 3NF
 - ↳ For every non-trivial FD, $(X \rightarrow Y) \notin F^+$, X contains a key
 - ↳ For any R in the schema
 - IF (non-trivial $X \rightarrow Y$ holds and X does not contain a key)
 - Compute X^+
 - Decompose R into $(R_1 \text{ and } R_2)$
 - ↳ X is common
 - ↳ Y is all attr. in R except X
 - Repeat
 - ↳ Lossless, not unique

MapReduce

- Ex)

```
lines = sc.textFile("input.txt")
words = lines.flatMap(lambda line: line.split(" "))
wordCounts = words.map(lambda word: (word, 1))
wordCounts = wordCounts.reduceByKey(lambda a, b: a+b)
wordCounts.saveAsTextFile("output")
```

Disk

- Access Time = Seek Time + Rotational Delay + Transfer Time
 - ↳ Seek Time: t to move disk between tracks
 - ↳ RD: depends on spin speed
 - ↳ Avg. 1 full \odot and 0
 - ↳ TT: time to read track
 - ↳ RPM and sectors/track
- Transfer rate: rate at which data is transferred
 - ↳ $\frac{\text{RPM}}{60} \times \frac{\text{Sectors}}{\text{Track}} \times \frac{\text{bytes}}{\text{Sector}} = \text{burst}$
- Random I/O
 - ↳ More expensive for mag disk
- Platter, track, cylinder, block/sector

MongoDB

• Insertion: insertX(doc(s))

↳ db.books.insertOne({title: "a", attr: 1})

↳ db.books.insertMany({title: "a"}, {title: "b"})

• Retrieval: findX(cond)

↳ db.books.findOne({likes: 100})

↳ db.books.find({\$and: [{likes: {\$gte: 100}}, {likes: {\$lte: 200}}]})

• Update: updateX(cond, update)

↳ update = { \$set: { title: "MongoDB" } }

• Deletion: deleteX(cond)

• Agg.

↳ aggregate([pipeline])

↳ db.orders.aggregate([

{ \$match: { status: "A" } },

{ \$group: {

_id: "\$cust_id",

total: { \$sum: "\$amount" },

count: { \$sum: 1 } }

]})

↳ \$sort: { \$sort: { total: -1 } }

↳ \$limit: 10, -1 is DESC

↳ \$limit: { \$limit: { positive: int } }

↳ \$project: { \$project: { _id: 0, 'gr': 1 } }

↳ \$unwind: { \$unwind: "\$" }

Files

• Spanned vs. unspanned

↳ Spanned: tuples cross blocks

↳ unspanned: tuples within 1 block

↳ More common

↳ wastes up to ~50% of block

• Variable-len tuples

↳ Reserve space - max space each tuple

↳ no max waste

↳ Variable space: tight packing

↳ fragmentation or insert and reshuffle

↳ Slotted page: array of pointers point

to tuples within blocks

• Index

↳ primary/clustering: table organized by key

↳ dense: 1 index entry per tuple

↳ sparse: 1 index entry per block

↳ secondary/non-clustering: unordered

↳ dense 1st level index

↳ Overflow => more disk I/Os over time

↳ B+AM not suitable for dynamic environment

B+ tree

↳ n = ptr spaces/node

↳ leaf node

↳ all pointers (except last) point to tuples

↳ 1/2 spaces used

↳ Non-leaf

↳ 1/2 spaces used

↳ Root

↳ 2 spaces used

↳ Usage

↳ Leaf: $\lceil \frac{n+1}{2} \rceil$ ptrs, $\lceil \frac{n+1}{2} \rceil - 1$ keys

↳ Non-leaf: $\lceil \frac{n}{2} \rceil$ ptrs, $\lceil \frac{n}{2} \rceil - 1$ keys

↳ Root: 2 ptrs, 1 key

↳ Insertion

↳ Leaf

↳ split node in 2

↳ split key between nodes

↳ 1st key of new node copied to parent + ptr to new node

↳ Non-leaf

↳ split in 2, max middle key to parent

↳ new root

↳ allocate new node

↳ Deletion

↳ Leaf

↳ Coalesce

↳ can merge w/ neighbor

↳ update ptrs, delete key

↳ Redistribute

↳ fill node from siblings

↳ update ptr in parent

↳ Non-leaf

↳ Coalesce

↳ merge w/ neighbor by pulling splitting key from parent

↳ delete pointer to merged node

↳ Redistribute

↳ overflow node by pulling splitting key down

↳ BPP overflow also

↳ Root

↳ delete empty node

↳ n determined by node size, search key size, ptr size

Hash index

↳ Extendible hashing: use 1st 3 bits from hash value. if needed to point queries, not range

↳ Insertion

↳ start @ i=0

↳ overflow

↳ i = direction

↳ double dir'n size by copying ptrs

↳ directory bit

↳ new tuple from overflowing bucket to new

↳ update directory ptr

↳ i++

- Files
 - Hash index
 - ↳ deletion
 - ↳ merge cond.
 - ↳ Bucket
 - ↳ bucket is are same
 - ↳ First (1-1) bits to hash key vals are same
 - ↳ Dir.
 - ↳ all bucket is < dir. i

- Join
 - NLJ
 - ↳ for $r \in R$ for $s \in S$
 - if $r.A = s.A \Rightarrow (r, s)$
 - IS
 - ↳ create index for $S.A$ for $r \in R$
 - $x = \text{lookup index on } S.A \text{ for } r.A$
 - for $s \in S$
 - if $s = x$

- SMJ
 - ↳ sort R/S by A
 - $i = j = 1$
 - while $(i \leq |R| \text{ or } j \leq |S|)$
 - if $(R[i].A = S[j].A)$
 - $(R[i], S[j])$
 - $i++, j++$
 - elif $(R[i].A > S[j].A)$ $j++$
 - else $i++$

- HTJ
 - ↳ hash R tuples into $C_1 \dots C_k$ buckets
 - hash S tuples into $H_1 \dots H_k$ buckets
 - for $i = 1$ to k
 - match tuples in C_i, H_i buckets

- SMJ
 - ↳ Read from R into m
 - ↳ sort R subset
 - ↳ Generate partitions of data
 - ↳ Read blocks from sorted runs into m ($m-1$)
 - ↳ Place smallest vals into staging block \rightarrow
 - ↳ init, sort, + $\lceil \log_{m-1}(\frac{n}{m}) \rceil$ merges
 - ↳ sort: $2bp(\lceil \log_{m-1}(\frac{n}{m}) \rceil + 1)$
 - ↳ $3m = bp + bs$
 - ↳ $2bp(\lceil \log_{m-1}(\frac{n}{m}) \rceil + 1) + 2bs(\lceil \log_{m-1}(\frac{n}{m}) \rceil + 1) + (bp + bs)$

- NLJ
 - ↳ Block first runs a block for S
 - ↳ B/As as many R into m as poss.
 - ↳ Scan S once for each group of R
 - ↳ $bp + (br/m - 2)bs$

- Transactions
 - COMMIT/ROLLBACK
 - ACID
 - ↳ Atomicity: all or none
 - ↳ Consistency: consistent before, consistent after
 - ↳ Isolation: concurrent results = sequential results
 - ↳ Durability: stable after crash

- Isolation lvs
 - ↳ Dirty Read - reading val from uncommitted T
 - ↳ READ UNCOMMITTED
 - ↳ Non-Repeatable Read - T_i gets diff values on multiple reads
 - ↳ READ COMMITTED
 - ↳ Phantom - statements see parts of new transactions
 - ↳ REPEATABLE READ
 - ↳ SERIALIZABLE
 - ↳ READ ONLY = Only SELECT
 - ↳ SET TRANSACTION [access_mode] ISOLATION LEVEL <level>
 - ↳ Eye of the beholding operation
 - Loggings
 - ↳ $\langle T_i, \text{start} \rangle$
 - $\langle T_i, \text{commit/abort} \rangle$
 - $\langle T_i, x, \text{old-value}, \text{new-value} \rangle$
 - ↳ Before commit, all records must flush to disk first
 - ↳ After tuple written back to disk, all log rec-run flush to disk first
 - ↳ $\langle T_i, A, S, 10 \rangle$ to disk before $M:10$
 - ↳ ROLLBACK \Rightarrow DBMS reverts to old value
 - ↳ Crash, DBMS re-executes log commands then rolls back all non-committed actions