

# CS143

# Advanced SQL

Professor Junghoo “John” Cho

# What to Learn

- Aggregate functions
- Window function
- Case expression
- Order by and Fetch first
- Data modification
- NULL and three-valued logic
- Outer join
- Multiset semantic for set operators
- SQL expressive power and recursion

# Q1: Average GPA of all students

Student(sid, name, addr, age, GPA)

| sid | name   | addr         | age | GPA |
|-----|--------|--------------|-----|-----|
| 301 | Andy   | 183 Westwood | 19  | 2.1 |
| 303 | Elaine | 301 Wilshire | 17  | 3.9 |
| 401 | James  | 183 Westwood | 17  | 3.5 |
| 208 | Esther | 421 Wilshire | 20  | 3.1 |

Class(dept, cnum, sec, unit, title, instructor)

| dept | cnum | sec | unit | title      | instructor   |
|------|------|-----|------|------------|--------------|
| CS   | 112  | 01  | 03   | Modeling   | Dick Muntz   |
| CS   | 143  | 01  | 04   | DB Systems | John Cho     |
| EE   | 143  | 01  | 03   | Signal     | Dick Muntz   |
| ME   | 183  | 02  | 05   | Mechanics  | Susan Tracey |

Enroll(sid, dept, cnum, sec)

| sid | dept | cnum | sec |
|-----|------|------|-----|
| 301 | CS   | 112  | 01  |
| 301 | CS   | 143  | 01  |
| 303 | EE   | 143  | 01  |
| 303 | CS   | 112  | 01  |
| 401 | CS   | 112  | 01  |

# Key Challenge of Q1

- What we learned: Information from *one* input tuple per output

| sid            | name              | addr                    | age           | GPA            |
|----------------|-------------------|-------------------------|---------------|----------------|
| <del>301</del> | <del>Andy</del>   | <del>183 Westwood</del> | <del>19</del> | <del>2.1</del> |
| 303            | Elaine            | 301 Wilshire            | 17            | 3.9            |
| 401            | James             | 183 Westwood            | 17            | 3.5            |
| <del>208</del> | <del>Esther</del> | <del>421 Wilshire</del> | <del>20</del> | <del>3.1</del> |

GPA > 3.2 →

|     |
|-----|
| 3.9 |
| 3.5 |

- What we need to do: Combine information from *multiple input tuples* into a *single* output tuple

| sid | name   | addr         | age | GPA |
|-----|--------|--------------|-----|-----|
| 301 | Andy   | 183 Westwood | 19  | 2.1 |
| 303 | Elaine | 301 Wilshire | 17  | 3.9 |
| 401 | James  | 183 Westwood | 17  | 3.5 |
| 208 | Esther | 421 Wilshire | 20  | 3.1 |

AVG(GPA) →

|      |
|------|
| 3.15 |
|------|

# Q1: Average GPA of all students

Student(sid, name, addr, age, GPA)

| sid | name   | addr         | age | GPA |
|-----|--------|--------------|-----|-----|
| 301 | Andy   | 183 Westwood | 19  | 2.1 |
| 303 | Elaine | 301 Wilshire | 17  | 3.9 |
| 401 | James  | 183 Westwood | 17  | 3.5 |
| 208 | Esther | 421 Wilshire | 20  | 3.1 |

Class(dept, cnum, sec, unit, title, instructor)

| dept | cnum | sec | unit | title      | instructor   |
|------|------|-----|------|------------|--------------|
| CS   | 112  | 01  | 03   | Modeling   | Dick Muntz   |
| CS   | 143  | 01  | 04   | DB Systems | John Cho     |
| EE   | 143  | 01  | 03   | Signal     | Dick Muntz   |
| ME   | 183  | 02  | 05   | Mechanics  | Susan Tracey |

Enroll(sid, dept, cnum, sec)

| sid | dept | cnum | sec |
|-----|------|------|-----|
| 301 | CS   | 112  | 01  |
| 301 | CS   | 143  | 01  |
| 303 | EE   | 143  | 01  |
| 303 | CS   | 112  | 01  |
| 401 | CS   | 112  | 01  |

# Aggregate Functions

- Allows “aggregating” results from multiple tuples to produce a single output tuple
- AVG, SUM, COUNT, MIN, MAX on single attribute
  - COUNT(\*): counts the number of matching tuples

# Q2: Number of students taking CS classes

Student(sid, name, addr, age, GPA)

| sid | name   | addr         | age | GPA |
|-----|--------|--------------|-----|-----|
| 301 | Andy   | 183 Westwood | 19  | 2.1 |
| 303 | Elaine | 301 Wilshire | 17  | 3.9 |
| 401 | James  | 183 Westwood | 17  | 3.5 |
| 208 | Esther | 421 Wilshire | 20  | 3.1 |

Class(dept, cnum, sec, unit, title, instructor)

| dept | cnum | sec | unit | title      | instructor   |
|------|------|-----|------|------------|--------------|
| CS   | 112  | 01  | 03   | Modeling   | Dick Muntz   |
| CS   | 143  | 01  | 04   | DB Systems | John Cho     |
| EE   | 143  | 01  | 03   | Signal     | Dick Muntz   |
| ME   | 183  | 02  | 05   | Mechanics  | Susan Tracey |

Enroll(sid, dept, cnum, sec)

| sid | dept | cnum | sec |
|-----|------|------|-----|
| 301 | CS   | 112  | 01  |
| 301 | CS   | 143  | 01  |
| 303 | EE   | 143  | 01  |
| 303 | CS   | 112  | 01  |
| 401 | CS   | 112  | 01  |

# Q3: Average GPA of students who take CS classes

Student(sid, name, addr, age, GPA)

| sid | name   | addr         | age | GPA |
|-----|--------|--------------|-----|-----|
| 301 | Andy   | 183 Westwood | 19  | 2.1 |
| 303 | Elaine | 301 Wilshire | 17  | 3.9 |
| 401 | James  | 183 Westwood | 17  | 3.5 |
| 208 | Esther | 421 Wilshire | 20  | 3.1 |

Class(dept, cnum, sec, unit, title, instructor)

| dept | cnum | sec | unit | title      | instructor   |
|------|------|-----|------|------------|--------------|
| CS   | 112  | 01  | 03   | Modeling   | Dick Muntz   |
| CS   | 143  | 01  | 04   | DB Systems | John Cho     |
| EE   | 143  | 01  | 03   | Signal     | Dick Muntz   |
| ME   | 183  | 02  | 05   | Mechanics  | Susan Tracey |

Enroll(sid, dept, cnum, sec)

| sid | dept | cnum | sec |
|-----|------|------|-----|
| 301 | CS   | 112  | 01  |
| 301 | CS   | 143  | 01  |
| 303 | EE   | 143  | 01  |
| 303 | CS   | 112  | 01  |
| 401 | CS   | 112  | 01  |



# Q4: Average GPA for each age group

Student(sid, name, addr, age, GPA)

| sid | name   | addr         | age | GPA |
|-----|--------|--------------|-----|-----|
| 301 | Andy   | 183 Westwood | 19  | 2.1 |
| 303 | Elaine | 301 Wilshire | 17  | 3.9 |
| 401 | James  | 183 Westwood | 17  | 3.5 |
| 208 | Esther | 421 Wilshire | 20  | 3.1 |



| Age | AVG(GPA) |
|-----|----------|
| 17  | 3.7      |
| 19  | 2.1      |
| 20  | 3.1      |

# GROUP BY and SELECT attributes

- Q: Is the following query meaningful?

```
SELECT sid, age, AVG(GPA)
FROM Student
GROUP BY age;
```

- With GROUP BY, SELECT can have only aggregate functions or attributes that have a single value in each group

# Q5: Number of classes each student takes

Student(sid, name, addr, age, GPA)

| sid | name   | addr         | age | GPA |
|-----|--------|--------------|-----|-----|
| 301 | Andy   | 183 Westwood | 19  | 2.1 |
| 303 | Elaine | 301 Wilshire | 17  | 3.9 |
| 401 | James  | 183 Westwood | 17  | 3.5 |
| 208 | Esther | 421 Wilshire | 20  | 3.1 |

Class(dept, cnum, sec, unit, title, instructor)

| dept | cnum | sec | unit | title      | instructor   |
|------|------|-----|------|------------|--------------|
| CS   | 112  | 01  | 03   | Modeling   | Dick Muntz   |
| CS   | 143  | 01  | 04   | DB Systems | John Cho     |
| EE   | 143  | 01  | 03   | Signal     | Dick Muntz   |
| ME   | 183  | 02  | 05   | Mechanics  | Susan Tracey |

Enroll(sid, dept, cnum, sec)

| sid | dept | cnum | sec |
|-----|------|------|-----|
| 301 | CS   | 112  | 01  |
| 301 | CS   | 143  | 01  |
| 303 | EE   | 143  | 01  |
| 303 | CS   | 112  | 01  |
| 401 | CS   | 112  | 01  |

Q: What about students who take no classes

# Q6: Students who take two or more classes

Student(sid, name, addr, age, GPA)

| sid | name   | addr         | age | GPA |
|-----|--------|--------------|-----|-----|
| 301 | Andy   | 183 Westwood | 19  | 2.1 |
| 303 | Elaine | 301 Wilshire | 17  | 3.9 |
| 401 | James  | 183 Westwood | 17  | 3.5 |
| 208 | Esther | 421 Wilshire | 20  | 3.1 |

Class(dept, cnum, sec, unit, title, instructor)

| dept | cnum | sec | unit | title      | instructor   |
|------|------|-----|------|------------|--------------|
| CS   | 112  | 01  | 03   | Modeling   | Dick Muntz   |
| CS   | 143  | 01  | 04   | DB Systems | John Cho     |
| EE   | 143  | 01  | 03   | Signal     | Dick Muntz   |
| ME   | 183  | 02  | 05   | Mechanics  | Susan Tracey |

Enroll(sid, dept, cnum, sec)

| sid | dept | cnum | sec |
|-----|------|------|-----|
| 301 | CS   | 112  | 01  |
| 301 | CS   | 143  | 01  |
| 303 | EE   | 143  | 01  |
| 303 | CS   | 112  | 01  |
| 401 | CS   | 112  | 01  |

# HAVING Clause

- Check aggregate conditions
  - Example: Students who take two classes or more
- Appear after GROUP BY

# Next Topic

- Aggregate functions
- **Window function**
- Case expression
- Order by and Fetch first
- Data modification
- NULL and three-valued logic
- Outer join
- Multiset semantic for set operators
- SQL expressive power and recursion

Q7: Per each student, return their name, GPA and the overall GPA average

| sid | name   | addr         | age | GPA |
|-----|--------|--------------|-----|-----|
| 301 | Andy   | 183 Westwood | 19  | 2.1 |
| 303 | Elaine | 301 Wilshire | 17  | 3.9 |
| 401 | James  | 183 Westwood | 17  | 3.5 |
| 208 | Esther | 421 Wilshire | 20  | 3.1 |



| name   | GPA | AVG(GPA) |
|--------|-----|----------|
| Andy   | 2.1 | 3.15     |
| Elaine | 3.9 | 3.15     |
| James  | 3.5 | 3.15     |
| Esther | 3.1 | 3.15     |

- Q: Will this work?  
SELECT name, GPA, AVG(GPA)  
FROM Student;
- Correct answer: Use window function!  
SELECT name, GPA, AVG(GPA) **OVER()**  
FROM Student;


# Window Function

- Introduced in SQL 2003
- Syntax: FUNCTION(*attr*) OVER()
  - Use the same aggregate FUNCTION(*attr*), but append OVER()
  - Example: MAX(GPA) OVER()
- Interpretation
  - Generate *one output tuple per input tuple*, but FUNCTION(*attr*) is computed over *all* input tuples



Q8: Per each student, return their name, GPA and the average of GPA their age group

| sid | name   | addr         | age | GPA |
|-----|--------|--------------|-----|-----|
| 301 | Andy   | 183 Westwood | 19  | 2.1 |
| 303 | Elaine | 301 Wilshire | 17  | 3.9 |
| 401 | James  | 183 Westwood | 17  | 3.5 |
| 208 | Esther | 421 Wilshire | 20  | 3.1 |



| name   | GPA | AVG(GPA) |
|--------|-----|----------|
| Andy   | 2.1 | 2.1      |
| Elaine | 3.9 | 3.7      |
| James  | 3.5 | 3.7      |
| Esther | 3.1 | 3.1      |

- Apply AVG(GPA) only within their “group” or “partition”, not over the entire input tuples
- PARTITION BY
  - SELECT name, GPA, AVG(GPA) OVER(PARTITION BY age)  
FROM Student;
  - PARTITION BY for window function  $\cong$  GROUP BY for aggregate function
- Read textbook Sec 5.5 to learn more on Window function
  - ORDER BY, RANK(), NTILE(), window range...

# CASE Expression

- Limited version of “If then else”
  - Returns different values depending on conditions
- Syntax:

```
CASE
  WHEN <condition> THEN <expr>
  WHEN <condition> THEN <expr>
  ELSE <expr>
END
```
- Can be used anywhere a column name can be referenced
  - SELECT, WHERE, GROUP BY, ...

## Q9: Average GPA within child/adult group

| sid | name   | addr         | age | GPA |
|-----|--------|--------------|-----|-----|
| 301 | Andy   | 183 Westwood | 19  | 2.1 |
| 303 | Elaine | 301 Wilshire | 17  | 3.9 |
| 401 | James  | 183 Westwood | 17  | 3.5 |
| 208 | Esther | 421 Wilshire | 20  | 3.1 |



| AVG(GPA) |
|----------|
| 3.7      |
| 2.6      |

(child group)  
(adult group)

## Q9: Average GPA within child/adult group

| sid | name   | addr         | age | GPA |
|-----|--------|--------------|-----|-----|
| 301 | Andy   | 183 Westwood | 19  | 2.1 |
| 303 | Elaine | 301 Wilshire | 17  | 3.9 |
| 401 | James  | 183 Westwood | 17  | 3.5 |
| 208 | Esther | 421 Wilshire | 20  | 3.1 |



| age_group | AVG(GPA) |
|-----------|----------|
| child     | 3.7      |
| adult     | 2.6      |

- What if we want to show “child” and “adult” as part of output, not just the average?

## Q9: Average GPA within child/adult group

| sid | name   | addr         | age | GPA |
|-----|--------|--------------|-----|-----|
| 301 | Andy   | 183 Westwood | 19  | 2.1 |
| 303 | Elaine | 301 Wilshire | 17  | 3.9 |
| 401 | James  | 183 Westwood | 17  | 3.5 |
| 208 | Esther | 421 Wilshire | 20  | 3.1 |



| childGPA | adultGPA |
|----------|----------|
| 3.7      | 2.6      |

- What about this output?

# ORDER BY

- SQL is based on multiset semantics
  - Duplicates are allowed
  - Tuple order is ignored
- Still, for presentation purposes, it may be useful to order the result tuples by certain attribute(s)
  - Example: Order student tuples by GPA
- ```
SELECT sid, GPA  
FROM Student  
ORDER BY GPA DESC, sid ASC
```

  - Default is ASC if omitted
  - ORDER BY does not change SQL semantics. It is purely for presentation

# Q10: Top-3 students ordered by their GPA

- Sometimes we just want a few rows from the result. Is there a way to limit the result size?
- A: `SELECT * FROM Students  
ORDER BY GPA DESC  
FETCH FIRST 3 ROWS ONLY`
- `FETCH FIRST` Clause in SQL 2008
  - `[ OFFSET <num> ROWS ] FETCH FIRST <count> ROWS ONLY`
  - Skip the first <num> tuples and return the subsequent <count> rows
  - Unfortunately, this was standardized too late. Many variations are being used
    - MySQL: `LIMIT <count> OFFSET <num>`

# General SQL SELECT

- SELECT attributes, aggregates  
FROM relations  
WHERE conditions  
GROUP BY attributes  
HAVING aggregate condition  
ORDER BY attributes  
FETCH FIRST n ROWS ONLY
- SELECT appears first, but is the last clause to be “interpreted”



# Data Modification in SQL

- Insert tuple (301, 'CS', 201, 01) to Enroll table
- Populate Honors table with students of  $\text{GPA} > 3.7$
- Syntax: `INSERT INTO relation tuples;`

# Data Modification in SQL

- Delete all students who are not taking classes
- Syntax: DELETE FROM *relation* WHERE *condition*;
- Increase all CS course numbers by 100
- Syntax: UPDATE *relation*  
SET  $A1 = V1, \dots, An = Vn$   
WHERE *condition*;

# SQL: More Tricky Details

- NULL values
- Outer join
- Bag semantics for set operators
- Expressive power of SQL and recursion

# Dealing with NULL

- Q: What will be returned from the following query if GPA is NULL?  
SELECT name  
FROM Student  
WHERE GPA \* 100/4 > 90
- Q: What should be the result from GPA \* 100?
  - If input to an *arithmetic operator* is NULL, its output is NULL
- Q: What should be the result from NULL > 90?
  - Arithmetic comparison with NULL returns **Unknown**

# 3-Valued Logic

- SQL is based on ***three-valued logic***.
  - All conditions are evaluated to be: **True**, **False** or **Unknown**
  - SQL returns a tuple if the result from condition is **True**
    - **False** or **Unknown** tuples will not be returned
- SELECT name  
FROM Student  
WHERE GPA \* 100/4 > 90

# Truth Table of Three-valued Logic

Assume GPA is NULL and age is 17

- Q: GPA > 3.7 AND age > 18. What is the result of this condition?
- Q: GPA > 3.7 OR age > 18. What is the result of this condition?

| AND   | True  | False | Unknown |
|-------|-------|-------|---------|
| True  | True  | False |         |
| False | False | False |         |

| OR    | True | False | Unknown |
|-------|------|-------|---------|
| True  | True | True  |         |
| False | True | False |         |

# NULL and Aggregates

- Q: What should be the result for the following queries?

SELECT SUM(GPA) FROM Student

SELECT AVG(GPA) FROM Student

SELECT COUNT(GPA) FROM Student

SELECT COUNT(\*) FROM Student

| sid | GPA  |
|-----|------|
| 1   | 3.0  |
| 2   | 3.6  |
| 3   | 2.4  |
| 4   | NULL |

# NULL and Aggregates

- Aggregate functions ignore NULL values
  - Except COUNT(\*), which counts a NULL valued tuple as a “valid” tuple
  - Note that COUNT(attr) does ignore a NULL valued attr
- When an input to an aggregate function is empty (= no input tuples):
  - COUNT() returns 0
  - All others return NULL



# NULL and Set Operators

- Q: What should be  $\{2.4, 3.0, \text{NULL}\} \cup \{3.6, \text{NULL}\}$ ?
- NULL is treated like other regular values for set operators

# Checking NULL

- In case we need to explicitly check whether an attribute value is NULL, we can use “IS NULL” or “IS NOT NULL” operator
  - Note that “= NULL” or “<> NULL” does **not** work!
- COALESCE() function
  - Return first non-NULL value in the list
  - Example: COALESCE(phone, email, addr)

Q: number of classes each student takes. return 0-class students as well

Student(sid, name, addr, age, GPA)

| sid | name   | addr         | age | GPA |
|-----|--------|--------------|-----|-----|
| 301 | Andy   | 183 Westwood | 19  | 2.1 |
| 303 | Elaine | 301 Wilshire | 17  | 3.9 |
| 401 | James  | 183 Westwood | 17  | 3.5 |
| 208 | Esther | 421 Wilshire | 20  | 3.1 |

Class(dept, cnum, sec, unit, title, instructor)

| dept | cnum | sec | unit | title      | instructor   |
|------|------|-----|------|------------|--------------|
| CS   | 112  | 01  | 03   | Modeling   | Dick Muntz   |
| CS   | 143  | 01  | 04   | DB Systems | John Cho     |
| EE   | 143  | 01  | 03   | Signal     | Dick Muntz   |
| ME   | 183  | 02  | 05   | Mechanics  | Susan Tracey |

Enroll(sid, dept, cnum, sec)

| sid | dept | cnum | sec |
|-----|------|------|-----|
| 301 | CS   | 112  | 01  |
| 301 | CS   | 143  | 01  |
| 303 | EE   | 143  | 01  |
| 303 | CS   | 112  | 01  |
| 401 | CS   | 112  | 01  |

Outer join preserves dangling tuples

# Outer Join

Student

| sid | name   |
|-----|--------|
| 301 | John   |
| 303 | Elaine |

Enroll

| sid | cid   |
|-----|-------|
| 301 | CS143 |
| 401 | CS112 |

- Student LEFT OUTER JOIN Enroll ON Student.sid = Enroll.sid
- Student RIGHT OUTER JOIN Enroll ON Student.sid = Enroll.sid
- Student FULL OUTER JOIN Enroll ON Student.sid = Enroll.sid

# SQL and Multiset Semantics

- Multiset (= Bag)
  - A set with duplicate elements
  - Order of elements does not matter
  - $\{a, a, b, c\} = \{a, b, c, a\} \neq \{a, b, c\}$
- SQL is based on multiset semantics
  - We already learned how duplicates are generated and kept in SQL
  - Use DISTINCT to eliminate duplicates in the result
  - Exception: ***set operators are based on set semantics***

# Multiset Semantics for Set Operators

- To use bag semantics for set operators, use **ALL** keyword
  - UNION ALL, INTERSECT ALL, EXCEPT ALL
- Q:  $\{a, a, b\} \cup \{a, b, c\}$ ?
- Q:  $\{a, a, a, b, c\} \cap \{a, a, b\}$ ?
- Q:  $\{a, a, b, b\} - \{a, b, b, c\}$ ?

# Multiset Semantics and Equivalence Rules

Under multiset semantics:

- $R \cup S = S \cup R$ ?
- $R \cap S = S \cap R$ ?
- $R \cap (S \cup T) = (R \cap S) \cup (R \cap T)$ ?
- Not all set equivalence rules hold under multiset. Be careful!!!

# Expressive Power of SQL

- Q: Find all ancestors of Susan

Parent

| child | parent |
|-------|--------|
| Susan | John   |
| John  | James  |
| James | Elaine |
| ...   | ...    |

- Q: Find all cities reachable from A?

Reachable

| City 1 | Citi 2 |
|--------|--------|
| A      | B      |
| B      | D      |
| B      | E      |
| ...    | ...    |



# Expressive Power of SQL

- SQL is a very expressive language, but its expressive power is limited
  - SQL is not a “Turing-complete” language
- For example, the closure of a set cannot be computed using SQL92
  - Example: all ancestors, all reachable nodes
  - Support for recursion is needed to compute a closure
- SQL99 added support for recursion
  - ```
WITH RECURSIVE Ancestor(child, ancestor) AS (  
    (SELECT * FROM Parent)  
    UNION  
    (SELECT P.child, A.ancestor  
     FROM Parent P, Ancestor A  
     WHERE P.parent = A.child) )  
SELECT ancestor FROM Ancestor WHERE child='Susan';
```
  - Read textbook for more details on SQL99 Recursion

# What We Learned

- Aggregate function
- Window function
- Case expression
- Order by and Fetch first
- Data Modification
- NULL and three-valued logic
- Outer join
- Multiset semantic for set operators
- SQL expressive power and recursion