

CS143

Map Reduce (Spark)

Professor Junghoo “John” Cho

Distributed Analytics using Cluster


- Often, our data is non-relational (e.g., flat file) and huge
 - Billions of query logs
 - Billions of web pages
 - ...
- Q: Can we perform analytics on large data quickly using thousands of machines?

Example 1: Search Log Analysis

- Log of billions of queries. Count frequency of each query
 - Input query log:
cat,time,userid1,ip1,referrer1
dog,time,userid2,ip2,referrer2
...
 - Output query frequency:
cat 200000
dog 120000
...
- Q: How can we perform this task? How can we parallelize it?


Example 1: Search Log Analysis (1)

- Step 1: “Transform” each line of query log into (query, 1)



Example 1: Search Log Analysis (2)


- Step 2: Collect all tuples with the same query and “aggregate” them



- Q: How can we parallelize the two steps?

Example 1: Search Log Analysis (3)

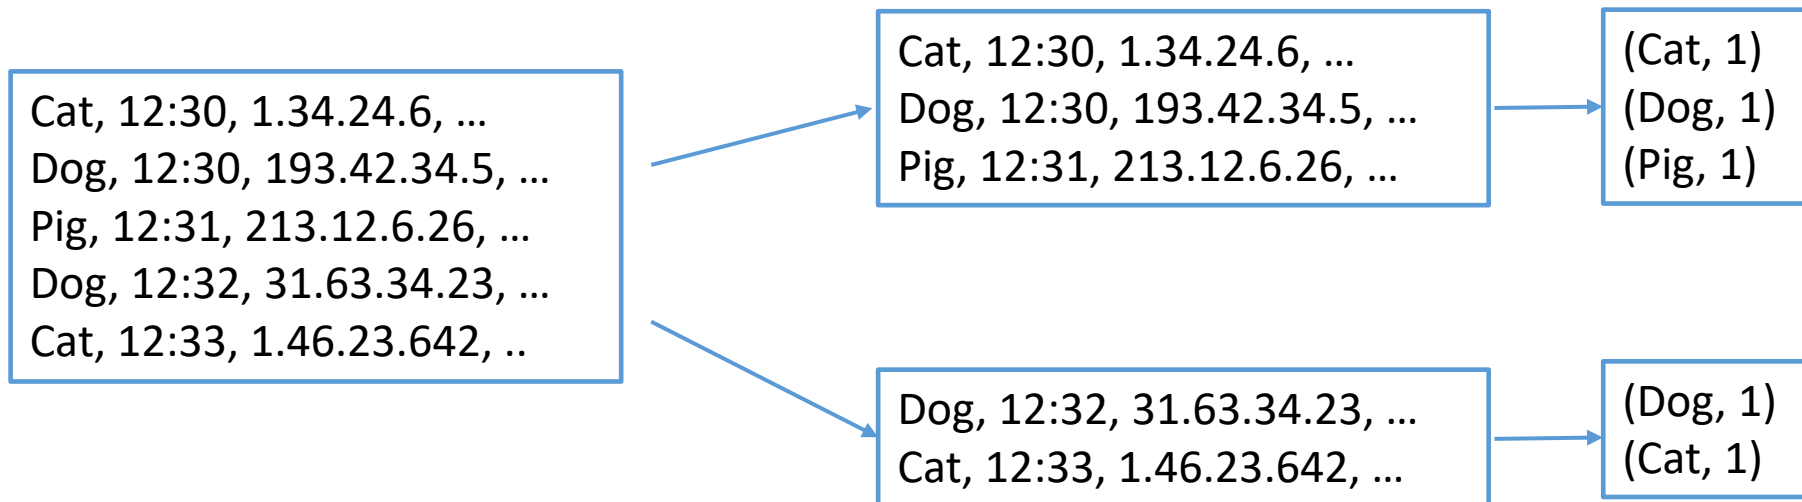
- Step 1: “Transform” each line of query log into (query, 1)



- Q: Can the transformation of each line be done independently of each other?

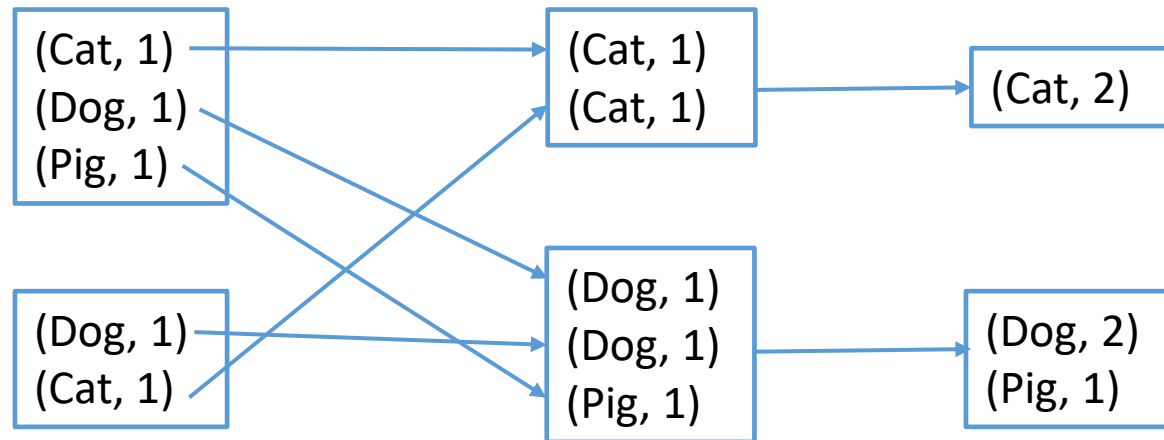
Example 1: Search Log Analysis (4)

- Step 1: For parallel processing
 - Split input data into multiple independent chunks
 - Move each chunk to separate machine
 - Perform “transformation” on multiple machines in parallel



Example 1: Search Log Analysis (5)

- Q: How do we parallelize the second “aggregation” step?
- Step 2: For parallel processing
 - Move the tuples with the same query to the same machine
 - Perform aggregation on multiple machine in parallel



Example 2: Web Indexing

- 1 billion pages. Build “inverted index”
 - Input documents:
 - 1: cat chases dog
 - 2: dog loves cat
 - ...
 - Output index:
 - cat 1,2,5,10,20
 - dog 1,2,3,8,9
- Q: How can we do this?

Example 2: Web Indexing (1)

- Step 1: “Transform” every document into (word, doc_id) tuples

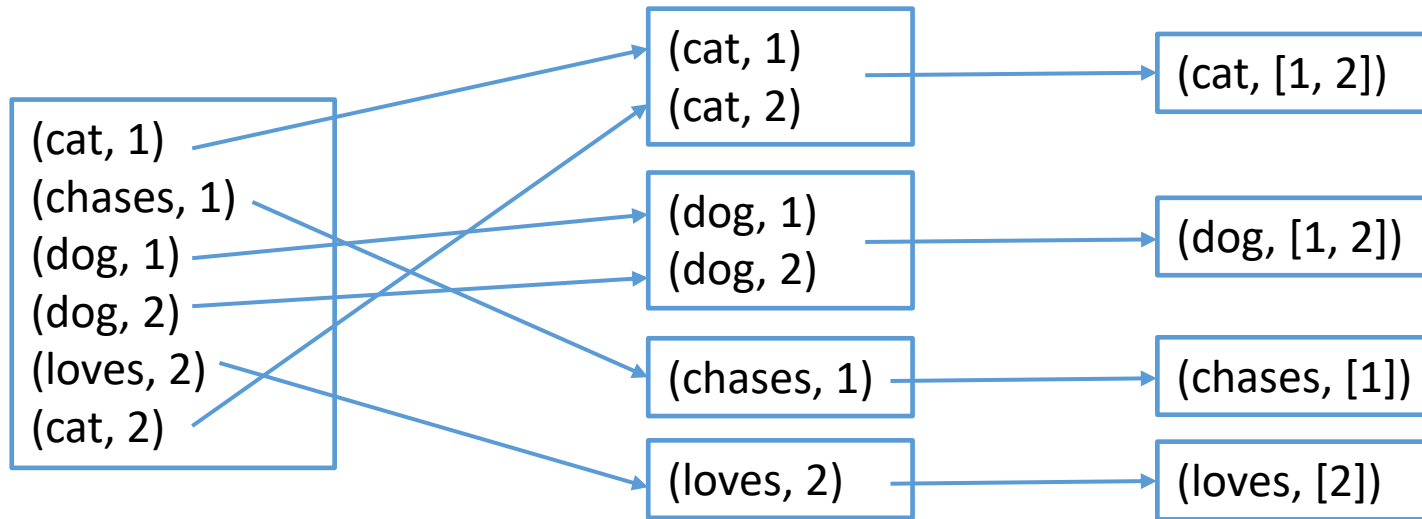
1: cat chases dog
2: dog loves cat



(cat, 1)
(chases, 1)
(dog, 1)
(dog, 2)
(loves, 2)
(cat, 2)

Example 2: Web Indexing (2)


- Step 2: Collect all tuples with the same word and “aggregate” (or concatenate) the doc_id's



- Q: How can we parallelize the two steps on multiple machines?

Example 2: Web Indexing (3)


- Step 1: “Transform” every document into (word, doc_id) tuples



- Q: Can the transformation of each document be done independently of each other?

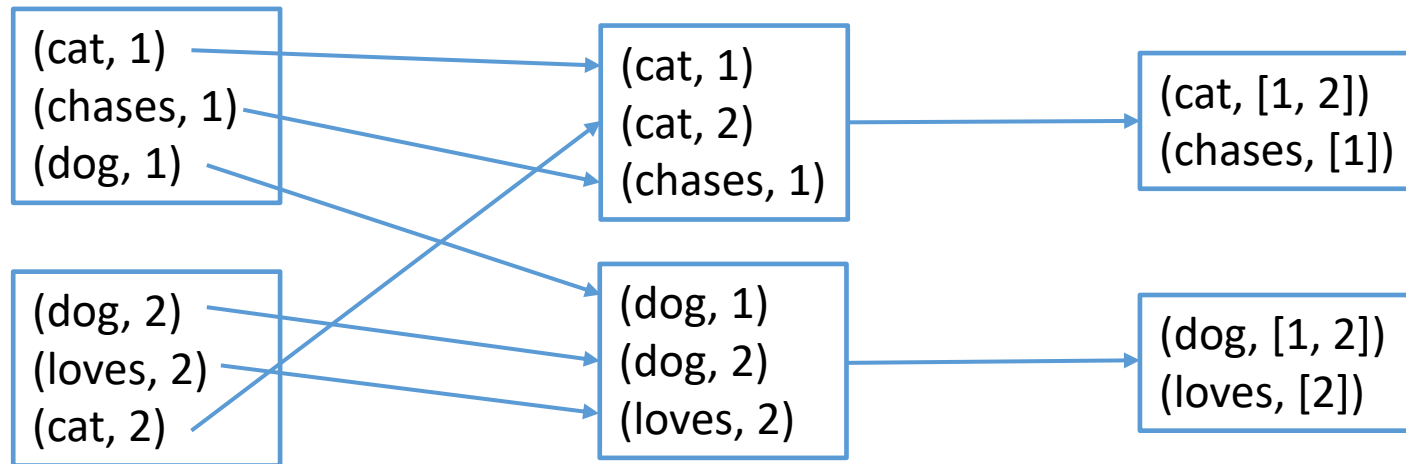
Example 2: Web Indexing (4)

- Step 1: For parallel processing
 - Split input data into multiple independent chunks
 - Move each chunk to separate machine
 - Perform “transformation” on multiple machines in parallel



Example 2: Web Indexing (5)

- Q: How can we parallelize second “concatenation step”?
- Step 2: For parallel processing
 - Move the tuples with the same word to the same machine
 - Perform aggregation on multiple machine in parallel



Generalization (1)

- Input data consists of multiple independent units
 - Each line of query log
 - Each web page
- Partition input data into multiple “chunks” and distribute them to multiple machines
- Transformation/map input into (key, value) tuples
 - Query log: $\text{query_log_line} \rightarrow (\text{query}, 1)$
 - Indexing: $\text{web_page} \rightarrow (\text{word}_1, \text{page_id}), (\text{word}_2, \text{page_id}), \dots$
- Reshuffle tuples of the same key to the same machine
- Aggregate/reduce the tuples of same keys
 - Query log: $(\text{query}, 1), (\text{query}, 1), \dots \rightarrow (\text{query}, \text{count})$
 - Indexing: $(\text{word}, 1), (\text{word}, 3), \dots \rightarrow (\text{word}, [1, 3, \dots])$
- Collect and output the aggregation results

Generalization (2)

- The two examples are almost the same except
 - “The mapping function”
 - Query log: $\text{query_log_line} \rightarrow (\text{query}, 1)$
 - Indexing: $\text{web_page} \rightarrow (\text{word}_1, \text{page_id}), (\text{word}_2, \text{page_id}), \dots$
 - “The reduction function”
 - Query log: $(\text{query}, 1), (\text{query}, 1), \dots \rightarrow (\text{query}, \text{count})$
 - Indexing: $(\text{word}, 1), (\text{word}, 3), \dots \rightarrow (\text{word}, [1, 3, \dots])$

MapReduce Model

- Programmer provides
 1. Map function: “unit data” $\rightarrow (k', v'), (k'', v''), \dots$
 2. Reduce function: $(k, v_1), (k, v_2), \dots \rightarrow (k, \text{aggr}(v_1, v_2, \dots))$
- MapReduce handles the rest
 - Automatic data partition, distribution, and collection
 - Failure and speed-disparity handling
- Many systems exist supporting MapReduce model

Hadoop

- First open-source implementation of MapReduce and GFS (Google File System)
 - Implemented in Java
- User implements map and reduce functions as:
 - `Mapper.map(key, value, output, reporter)`
 - `Reducer.reduce(key, value, output, reporter)`

Spark

- Open-source cluster computing infrastructure
- Supports MapReduce and SQL
 - Supports data flow more general than simple MapReduce
- Input data is converted into RDD (resilient distributed dataset)
 - A collection of independent tuples
 - The tuples are automatically distributed and shuffled by Spark
- Supports multiple programming languages
 - Scala, Java, Python, ...
 - Scala and Java are much more performant than others

Spark Example: Count words

```
lines = sc.textFile("input.txt")
words = lines.flatMap(lambda line: line.split(" "))
word1s = words.map(lambda word: (word, 1))
wordCounts = word1s.reduceByKey(lambda a,b: a+b)
wordCounts.saveAsTextFile("output")
```

Key Spark Functions

- Transformation: Convert RDD tuple into RDD tuple(s)
 - map(): convert one input tuple into one output tuple
 - flatMap(): convert one input into multiple output tuples
 - reduceByKey(): specify how two input “values” should be aggregated
 - filter(): filter out tuples based on condition
- Action: Perform “actions” on RDD
 - saveAsTextFile(): save RDD in a directory as text file(s)
 - collect(): create Python tuples from Spark RDD
 - textFile(): create RDD from text (each line becomes an RDD tuple)

What We Learned

- Large-scale data analytics on distributed cluster
- MapReduce model
- Spark