


# CS143: Files

Professor Junghoo “John” Cho

# Files: Main Problem



- How to store tables into disks?
- Q: 512Byte block. 80Byte tuple. How to store?

Jane	CS	3.7
Susan	ME	1.8
June	EE	2.6
Tony	CS	3.1





# Spanned vs Unspanned

- Unspanned




- Spanned



- Q: Maximum space waste for unspanned?


# Variable-Length Tuples

- How do we store them?



# Reserved Space

- Reserve the maximum space for each tuple




- Q: Any problem?

# Variable-Length Space

R1	R2	
R3		
R4	R5	
	R6	

- Pack tuples tightly
- Q: How do we know the end of a tuple?
- Q: What to do for delete/update?
- Q: How can we “point to” to a tuple?

# Slotted Page



Q: How can we point to a tuple?

# Long Tuples


- ProductReview(  
    pid INT,  
    reviewer VARCHAR(50),  
    date DATE,  
    rating INT,  
    comments VARCHAR(4000))
- Block size 512B
- How should we store it?




# Long Tuples

- Splitting tuples
  - Long attributes are stored separately (often as a separate file)

Block with short attributes.



Block with long attrs.



This block may also  
have fixed-length slots.

# Column-Oriented Storage

- `SELECT name FROM Students WHERE GPA > 3.7`
- For analytical queries, reading the entire row of a tuple may not be needed
  - Row-oriented storage forces us to read the entire row even if most columns are not needed for query processing

Elaine	1 Le Conte	3.7
James	3 Mississippi	2.8
John	12 Wilshire	1.8
Peter	4 Olympic	3.9
Susan	7 Pico	1.0
Tony	12 Sunset	2.4

# Column-Oriented Storage

- Store by column, not by row
- Unneeded Columns can be skipped for query processing
  - Better compression and caching behavior
- But
  - Column values of matching rows must be “joined”
  - Insertion/update of a row is more expensive (multiple IOs per row)

Elaine
James
John
Peter
Susan
Tony

1 Le Conte
3 Mississippi
12 Wilshire
4 Olympic
7 Pico
12 Sunset

3.7
2.8
1.8
3.9
1.0
2.4

# Sequential File


- Tuples are ordered by certain attribute(s) (search key)

Elaine	1 Le Conte	3.7
James	3 Mississippi	2.8
John	12 Wilshire	1.8
Peter	4 Olympic	3.9
Susan	7 Pico	1.0
Tony	12 Sunset	2.4

Search key: Name

# Sequencing Tuples

- Inserting a new tuple
  - Easy case




# Two Options

## 1) Rearrange

T1
T3
T6
T7
T8


## 2) Linked list



# Sequencing Tuples


- Inserting a new tuple
  - Difficult case

T1
T4
T5
T8
T9



# Sequencing Tuples

- Overflow page



- Reserving free space to avoid overflow
  - PCTFREE in DBMS  
CREATE TABLE R(a int) PCTFREE 40



# Things to Remember

- Spanned/unspanned tuples
- Variable-length tuples (slotted page)
- Long tuples
- Row-oriented vs column-oriented storage
- Sequential file and search key
  - Problems with insertion (overflow page)
  - PCTFREE

# CS143: Index

Professor Junghoo “John” Cho

# Topics to Learn

- Index
- Dense index vs. sparse index
- Primary index vs. secondary index  
(= clustering index vs. non-clustering index)
- Multi-level index
- Indexed Sequential Access Method (ISAM)

# Basic Problem

- `SELECT *`  
`FROM Student`  
`WHERE sid = 30`

sid	name	GPA
20	Susan	3.5
60	James	1.7
70	Peter	2.6
40	Elaine	3.9
30	Christy	2.9

- How can we answer the query?

# Random-Order File

- How do we find sid=30?

sid	name	GPA
20	Susan	3.5
60	James	1.7
70	Peter	2.6
40	Elaine	3.9
30	Christy	2.9

# Sequential File


- Table sequenced by sid. Find sid=30?

sid	name	GPA
20	Susan	3.5
30	James	1.7
40	Peter	2.6
50	Elaine	3.9
60	Christy	2.9



# Index: Basic Idea


- Build an “index” on the table
  - An auxiliary structure to help us quickly locate a tuple given a “search key”





# Dense, Primary Index

- Primary index (=clustering index)
  - Underlying table is sequenced by a key
  - Index is built on on the same key (= search key)
- Dense index
  - One (key, pointer) index entry per every tuple
- Search algorithm
  - Find the key from index and follow pointer
  - Maybe through binary search
- Q: Why dense index?
  - Isn't binary search on the file the same?




# Why Dense Index?

- Example
  - 100,000,000 tuples (900-bytes/tuple)
  - 4-byte search key, 4-byte pointer
  - 4096-byte block. Unspanned tuples
- Q: How many blocks for table (how big)?
- Q: How many blocks for index (how big)?

# Sparse, Primary Index

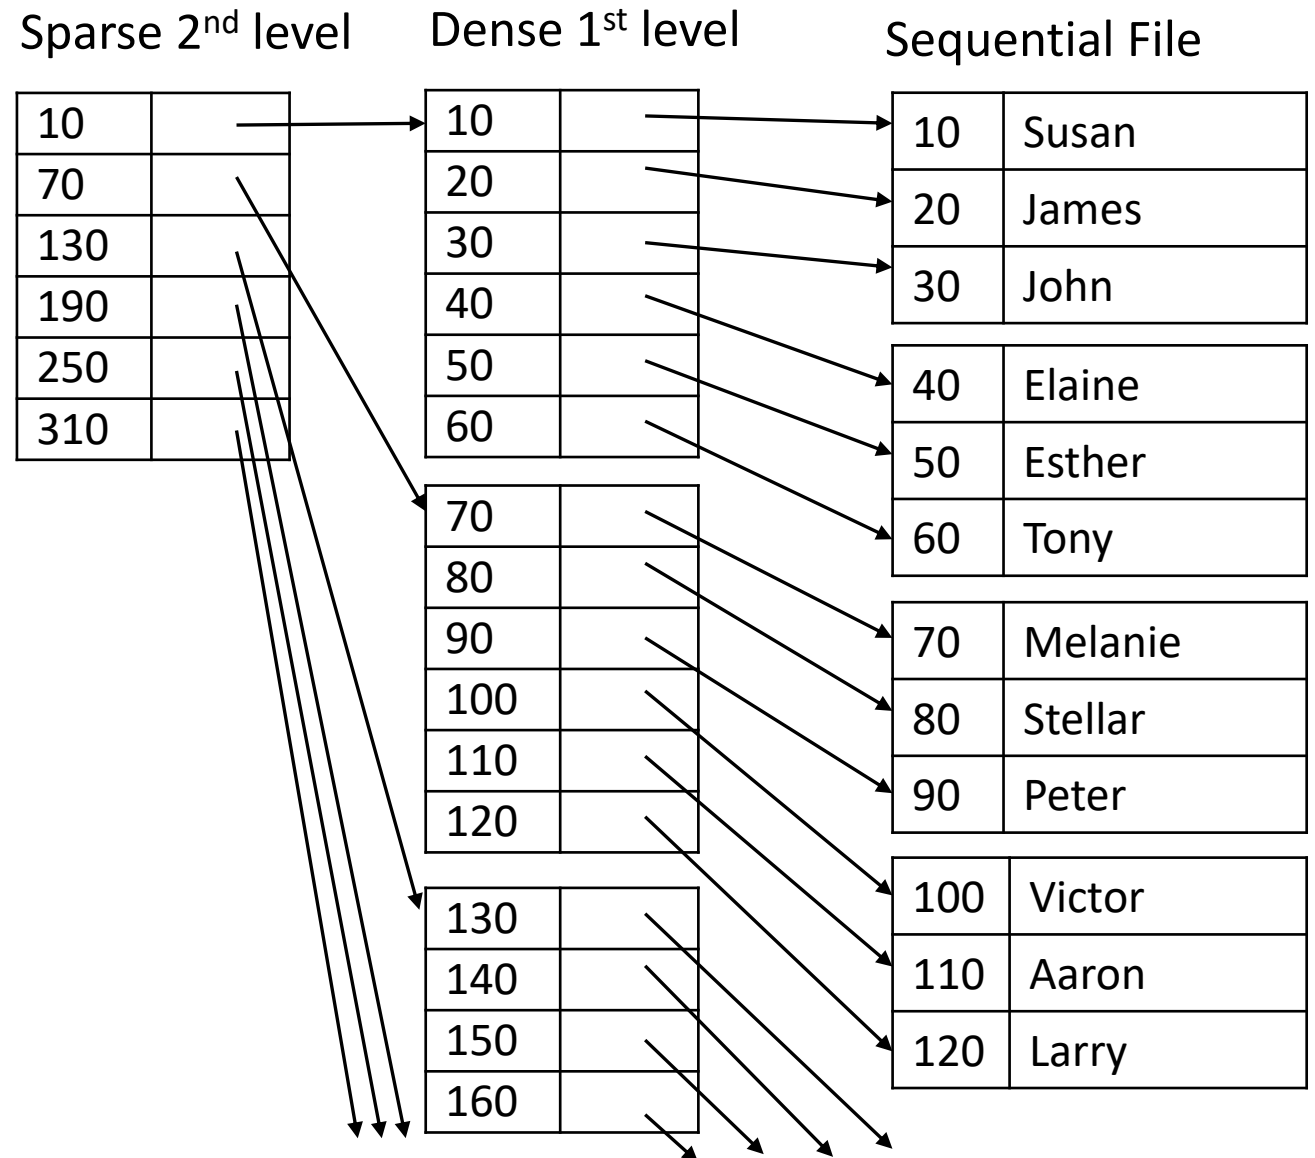
- Primary Index
  - Index is built on on the same search key as the underlying sequential file
- Sparse index
  - (key, pointer) pair per every “block”
  - (key, pointer) pair points to the first tuple in the block
- Q: How can we find 80?



# Multi-level index

Q: Why multi-level index?

Q: Does dense, 2nd level index make sense?



# Secondary (non-clustering) Index

- Secondary (non-clustering) index
  - When tuples in the table are not ordered by the index search key
    - Index on a non-search-key for sequential file
    - Unordered file
- Q: What index?
  - Does sparse index make sense?


Sparse Index

The diagram illustrates a sparse index. On the left, a 'Sparse Index' table contains index values: 110, 20, 70, 10, ..., ... Each index value has an arrow pointing to a specific tuple in a data table on the right. The data table contains tuples with a primary key and a name. The arrows show that the index values 110, 20, 70, and 10 point to the tuples for Aaron, James, Melanie, and Susan respectively. The other index values (20, 70, 10, ...) do not have arrows pointing to any tuple, demonstrating that a sparse index only contains a subset of the possible index values.

110		110	Aaron
20		40	Elaine
70		50	Esther
10		20	James
...		30	John
...		120	Larry
		70	Melanie
		90	Peter
		80	Stellar
		10	Susan
		60	Tony
		100	Victor

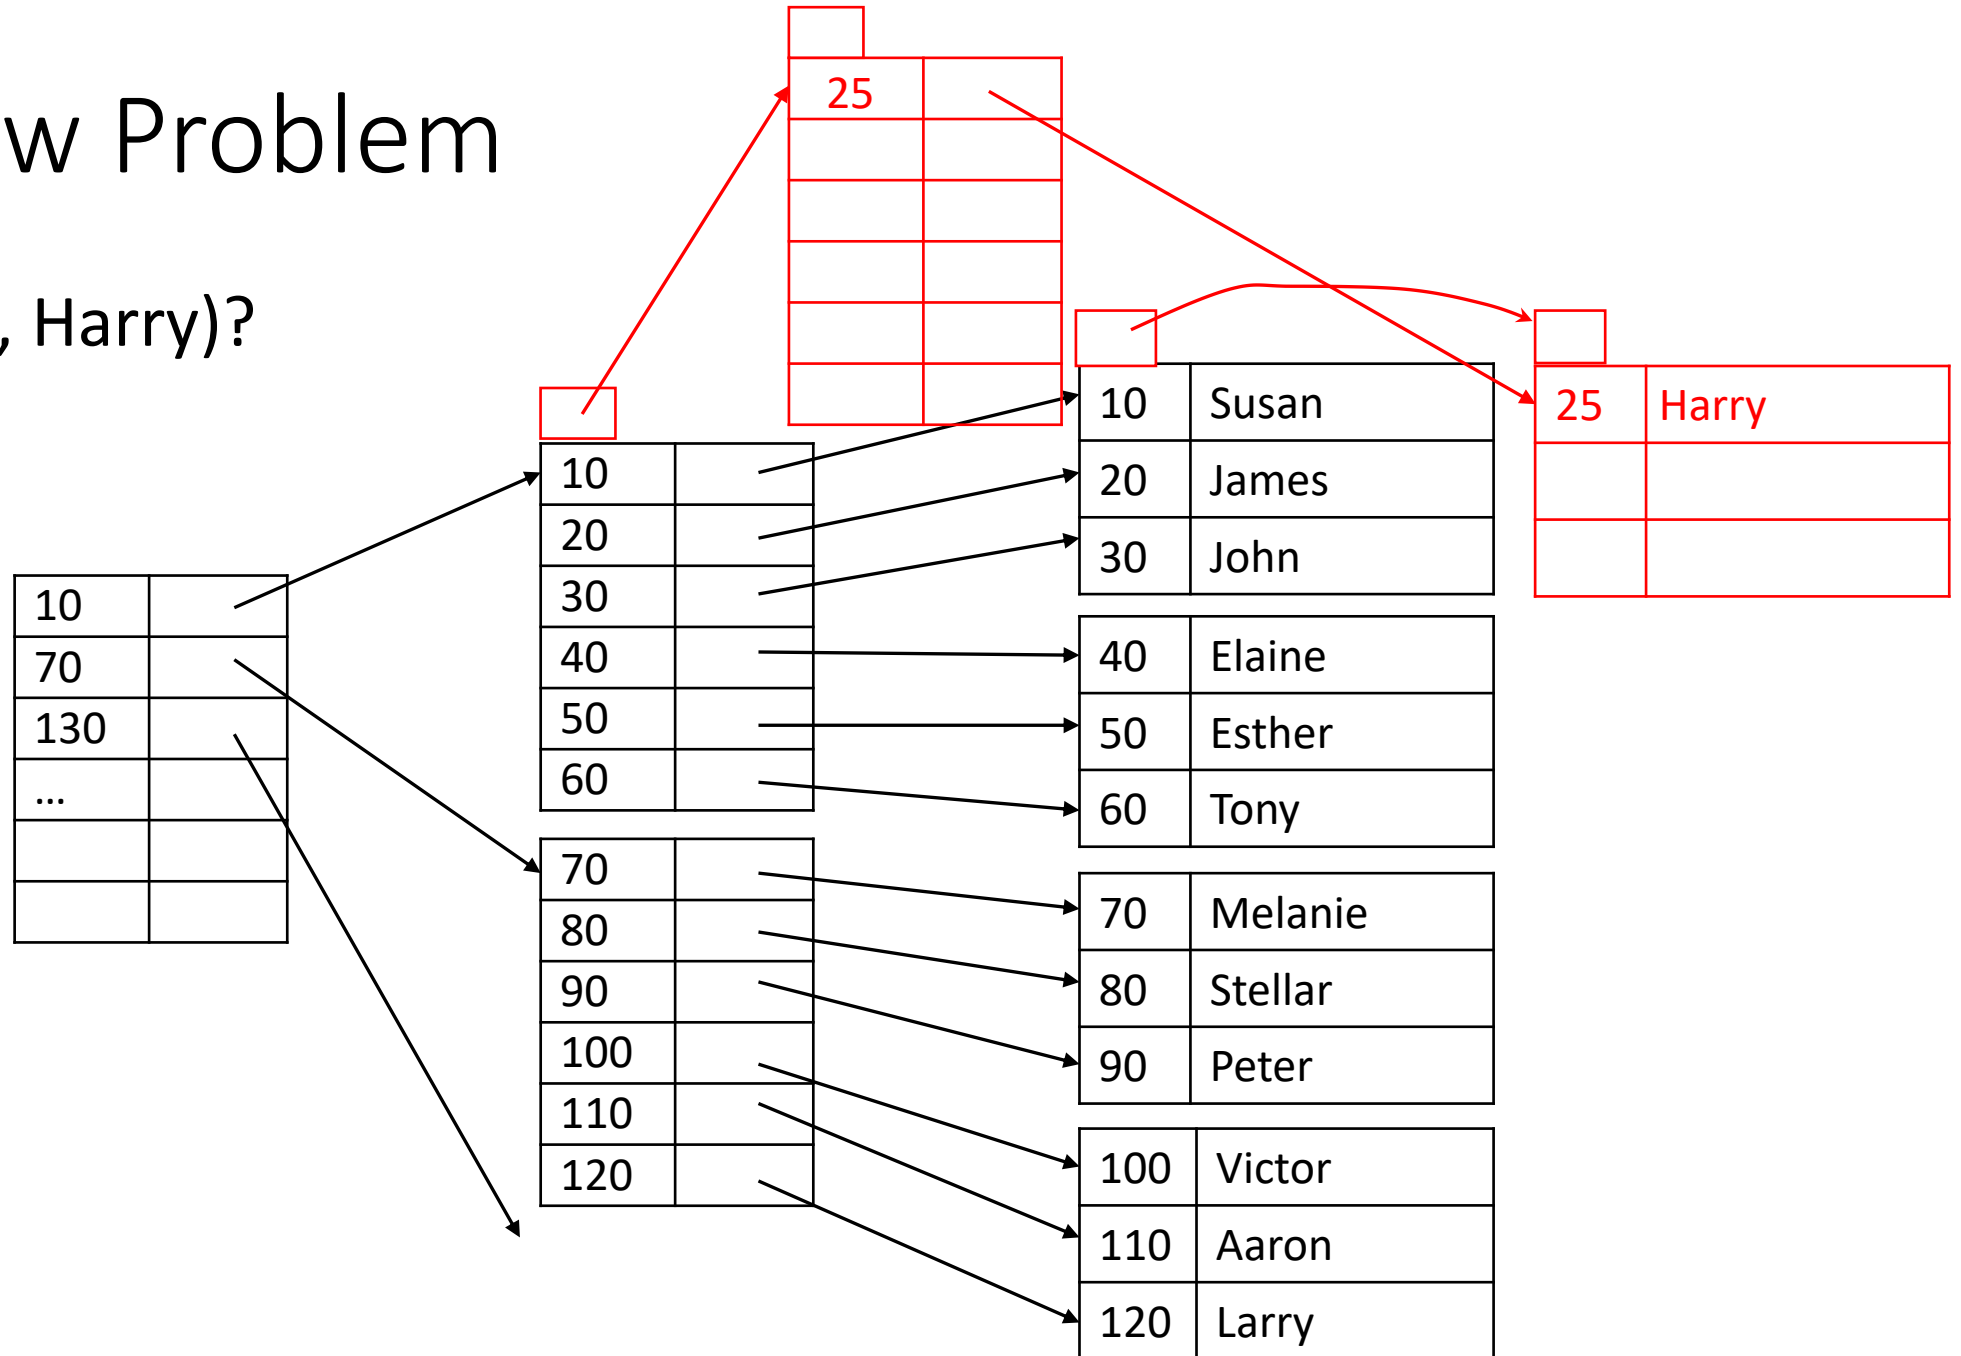
# Secondary index

- First level must be ***always dense***
- Sparse from second level



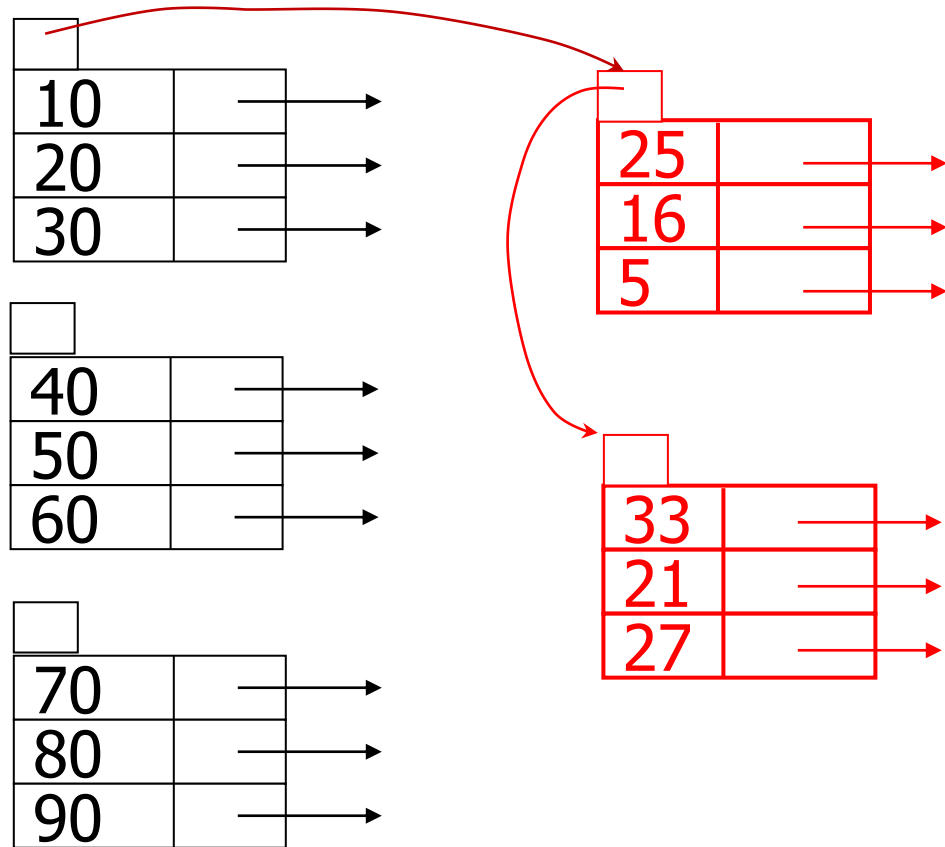
# Overflow Problem

Q: Insert (25, Harry)?



# Performance Problem after many insertions

- After many insertions, long chain of overflow pages





# Indexed Sequential Access Method (ISAM)

- Advantage
  - Simple
  - Sequential blocks
- Disadvantage
  - Not suitable for updates
  - Becomes ugly (loses sequentiality and balance) over time

# Index Creation in SQL

- CREATE INDEX <indexname> ON <table>(<attr>,<attr>,...)
- Example
  - CREATE INDEX sid\_idx ON Student(sid)
    - Creates a B+tree on the attributes
    - Speeds up lookup on sid

# Primary (Clustering) Index

- MySQL:
  - Primary key becomes the clustering index
- DB2:
  - CREATE INDEX idx ON Student(sid) CLUSTER
  - Tuples in the table are sequenced by sid
- Oracle: Index-Organized Table (IOT)
  - CREATE TABLE T (  
    ...  
    ) ORGANIZATION INDEX
  - B+tree on primary key
  - Tuples are stored at the leaf nodes of B+tree
- Periodic reorganization may still be necessary to improve range scan performance

# Important terms

- Search key (  $\neq$  primary key)
- Primary index vs. secondary index
  - Clustering index vs. non-clustering index
- Dense index vs. sparse index
- Multi-level index
- Indexed Sequential Access Method (ISAM)