

CS143: Joins

Professor Junghoo “John” Cho

Motivation

- Q: How do we process
SELECT * FROM Student WHERE sid > 30?
- Q: How do we process
SELECT * FROM Student S, Enroll E WHERE S.sid = E.sid?

$R \bowtie S ?$

R	A	
	40	T1
	60	T2
	30	T3
	10	T4
	20	T5

S	A	
	10	T6
	60	T7
	40	T8
	20	T9

Four Join Algorithms

- Nested-Loop Join (NLJ)
- Index Join (IJ)
- Sort-Merge Join (SMJ)
- Hash Join (HJ)

Nested-Loop Join (NLJ)

For each $r \in R$:

For each $s \in S$:

if $r.A = s.A$, then output (r,s)

R

40	T1
60	T2
30	T3
10	T4
20	T5

S

10	T6
60	T7
40	T8
20	T9


Index Join (IJ)

(1) Create an index for S.A if needed

(2) For each $r \in R$:

$X := \text{lookup index on S.A with } r.A \text{ value}$

For each $s \in X$, output (r,s)



Sort-Merge Join (SMJ)

- Sort the relations first, then join

R

10	T4
20	T5
30	T3
40	T1
60	T2

S

10	T6
20	T9
40	T8
60	T7

Sort-Merge Join (SMJ)

- (1) if not, sort R and S by A
- (2) $i \leftarrow 1; j \leftarrow 1;$
while $(i \leq |R|) \wedge (j \leq |S|)$:
 if $(R[i].A = S[j].A)$ then output $(R[i], S[j]); i \leftarrow i+1; j \leftarrow j+1;$
 else if $(R[i].A > S[j].A)$ then $j \leftarrow j+1$
 else if $(R[i].A < S[j].A)$ then $i \leftarrow i+1$

R

10	T4
20	T5
30	T3
40	T1
60	T2

S


10	T6
20	T9
40	T8
60	T7

Hash Join (HJ)


- Hash function: $h(v) \rightarrow [1, k]$
- Q: Given $(r \in R)$ and $(s \in S)$, can r and s join if $h(r.A) \neq h(s.A)$?
- Main idea
 - Partition tuples in R and S based on hash values on join attributes
 - Perform “joins” only between partitions of the same hash value

Hash Join (HJ)

- $H(k) = k \bmod 3$



40	T1
60	T2
30	T3
10	T4
20	T5



10	T6
60	T7
40	T8
20	T9

Hash Join (HJ)


Hash function: $h(v) \rightarrow [1, k]$

(1) Hashing stage (bucketizing): hash tuples into buckets

- Hash R tuples into $G1, \dots, Gk$ buckets
- Hash S tuples into $H1, \dots, Hk$ buckets

(2) Join stage: join tuples in matching buckets

- For $i = 1$ to k do
 match tuples in G_i, H_i buckets



Comparison of Join Algorithms

- Q: Which algorithm is better?
 - Q: What does “better” mean?
- Ultimate bottom line: Which algorithm is the “fastest”?
 - Q: How does the system know which algorithm runs fast? Run all join algorithms and pick the fastest one?


Cost Model

- A model to estimate the performance of a join algorithm
 - Multiple cost models are possible depending on their sophistication
- Our cost model: ***# disk blocks that are read/written during join***
 - Not perfect: ignores random vs sequential IO difference, CPU cost, ...
 - But simple to analyze
 - And “good enough” to pick the best join algorithm
 - Cost of join is dominated by disk IO
 - Most join algorithms have similar disk access pattern
 - ***Our cost model ignores the last IO for writing the final result***
 - This cost is the same for all algorithms

Running Example


- Join two tables: $R \bowtie S$
- $|R| = 1,000$ tuples, $|S| = 10,000$ tuples
- $b_R = 100$ blocks, $b_S = 1,000$ blocks (10 tuples/block)
- M = main memory “cache” 22 disk blocks

Memory




22 blocks

10 tuples



R (100 blocks)



S (1000 blocks)

Cost of Join Algorithms

	Cost	Formula ($b_R < b_S$)
NLJ		
SMJ		
HJ		
IJ		

Sort-Merge Join (SMJ)

- (1) if not, sort R and S by A
- (2) $i \leftarrow 1; j \leftarrow 1;$
while $(i \leq |R|) \wedge (j \leq |S|)$:
 if $(R[i].A = S[j].A)$ then output $(R[i], S[j]); i \leftarrow i+1; j \leftarrow j+1;$
 else if $(R[i].A > S[j].A)$ then $j \leftarrow j+1$
 else if $(R[i].A < S[j].A)$ then $i \leftarrow i+1$


R

10	T4
20	T5
30	T3
40	T1
60	T2

S

10	T6
20	T9
40	T8
60	T7

Cost of Join Stage of Sort-Merge Join




Q: Ignoring the final write of output, how many disk blocks are read during join?

Q: We only used 3 memory blocks. Can we use the rest to make things better?

What About?

- Q: Will this lead to fewer disk block reads?



Cost of Join Algorithms


	Cost ($M=22$, $b_R=100$, $b_S=1000$)	Formula ($b_R < b_S$)
NLJ		
SMJ		
HJ		
IJ		

Nested-Loop Join (NLJ): $R \bowtie S$

For each $r \in R$:

For each $s \in S$:


if $r.A = s.A$, then output (r,s)



Scan S table once for every tuple of R

Nested Loop Join


- Scan S table once for every tuple of R



- Q: Can we do better?

Block Nested Loop Join


- Scan S table once for every **block** of R



- Q: Can we do even better? What is the maximum # of blocks that we can read in one batch from R?

Block Nested Loop Join


- Scan S table once for every **20 blocks** of R



- Q: What if we read S first?

Block Nested Loop Join

- Scan R table once for every 20 blocks of S



Cost of Join Algorithms


	Cost ($M=22$, $b_R=100$, $b_S=1000$)	Formula ($b_R < b_S$)
NLJ		
SMJ		
HJ		
IJ		

Nested Loop Join Summary


- Always use block nested loop join (not the naïve algorithm)
- Read as many blocks as possible for the left table in one iteration
- Use the smaller table on the left (i.e., outer loop)

Hash Join (HJ)

- Step (1): Hashing stage: $h(v) \rightarrow [1, k]$




- Step (2): Join stage



HJ: Bucketizing Stage


- Read R table and hash them into k buckets



- Q: Given $M=22$, what is the maximum k ?
- Q: How many disk IOs to bucketize R ?

HJ: Bucketizing Stage


- Read S table and hash them into k buckets



- Q: In general, what is the cost for bucketizing R and S?

HJ: Join Stage

- Join tuples in G_i with those in H_i




- Q: How can we join tuples in G_1 with H_1 ? How should we use memory?

Cost of Join Algorithms

	Cost ($M=22$, $b_R=100$, $b_S=1000$)	Formula ($b_R < b_S$)
NLJ		
SMJ		
HJ		
IJ		

HJ: Join Stage

- Q: What if R is large, say $b_R = 1000$, and $G_i > 20$?



- A: Exactly the same as standard join problem. Apply “hash join” algorithm to join H_1 and G_1
 - Apply “hash join” algorithm using a new hash function!

HJ: Recursive Partitioning

- Use a new hash function $h'(v) \rightarrow [1, k]$ to recursively partition G_i and H_i to even smaller partitions (until one of them fit in main memory)
- # of bucketizing steps needed for R : $\left\lceil \log_{M-1} \frac{b_R}{M-2} \right\rceil$
 - In each bucketing steps, we perform $2(b_R + b_S)$ disk IOs

Cost of Join Algorithms


	Cost ($M=22$, $b_R=100$, $b_S=1000$)	Formula ($b_R < b_S$)
NLJ		
SMJ		
HJ		
IJ		

Index Join (IJ): $R \bowtie S$

For each $r \in R$:

$X := \text{lookup index on } S.A \text{ with } r.A \text{ value}$


For each $s \in X$, output (r,s)



- Cost = IOs for (R scan + index look up + tuple read from S)


IJ Example (1)

- 15 blocks for index
 - 1 root 14 leaf
- On average, 1 matching S tuple per an R tuple
- Q: How many disk IOs? How should we use the memory?




IJ Example (1)

- Cost for R scan:
- Cost for Index look up:
- Cost for read matching S tuple:




IJ Example (2)

- 40 blocks for index
 - 1 root 39 leaf
- On average, 10 matching S tuple per an R tuple
- Q: How many disk IOs? How should we use the memory?



IJ Example (2)

- Cost for R scan:
- Cost for Index look up:
- Cost for read matching S tuple:



Cost of Join Algorithms


	Cost ($M=22$, $b_R=100$, $b_S=1000$)	Formula ($b_R < b_S$)
NLJ		
SMJ		
HJ		
IJ		

SMJ: Cost of Sorting

- Sort-Merge Join
 1. Sort stage: Sort R and S
 2. Join stage: Join sorted R and S
- Q: How many disk IOs during sort stage?

SMJ: Cost of Sorting


- Q: How can we sort R?



- Q: How many blocks can we sort in each batch?
 - Do we need to allocate one block for output?
- Q: How many sorted runs?

SMJ: Cost of Sorting


- Q: What to do with sorted runs?



- Q: How many disk IOs during the “merge step” of sort?
- Q: Total IOs for sorting R?

SMJ: Cost of Sorting


- Q: How can we sort S?



- Q: How many sorted runs are produced from S?

SMJ: Cost of Sorting

- Q: How many sorted runs can we merge at a time?



- Q: What to do with the produced sorted runs?

SMJ: Cost of Sorting

- Q: How many "merging" steps are needed to sort S ?
 - 1 initial sorting
 - 2 merging steps of sorted runs
 - 2,000 disk IO's per each sorting/merging step
 - 6,000 total disk IO's to sort S table
- In general, to sort R of b_R blocks with M memory buffers, we need
 - 1 initial sorting
 - $\left\lceil \log_{M-1} \left(\frac{b_R}{M} \right) \right\rceil$ subsequent merging stages
 - $2 b_R$ disk IO's per each sorting/merging stage
 - In total, $2b_R \left(\left\lceil \log_{M-1} \left(\frac{b_R}{M} \right) \right\rceil + 1 \right)$ disk IO's are needed

Cost of Join Algorithms

	Cost ($M=22$, $b_R=100$, $b_S=1000$)	Formula ($b_R < b_S$)
NLJ		
SMJ		
HJ		
IJ		

Cost of Join Algorithms

	Cost (M=22, $b_R = 100, b_S = 1000$)	Formula ($b_R < b_S$)
NLJ	5,100	$b_R + \left\lceil \frac{b_R}{M-2} \right\rceil b_S$
SMJ	7,500 (if unsorted) 1,100 (if sorted)	$2b_R \left(\left\lceil \log_{M-1} \left(\frac{b_R}{M} \right) \right\rceil + 1 \right) +$ $2b_R \left(\left\lceil \log_{M-1} \left(\frac{b_R}{M} \right) \right\rceil + 1 \right) + (b_R + b_S)$
HJ	3,300	$2(b_R + b_S) \left\lceil \log_{M-1} \frac{b_R}{M-2} \right\rceil + (b_R + b_S)$
IJ	1,115 – 10,640	$b_R + R (C + J)$ C: index lookup cost, J: # matching S tuples per R tuple

Summary of Joins

- Nested-loop join is OK for “small” relations (relative to memory size)
- Hash join is usually the best for equi-join
 - If tables have not been sorted and with no index
 - Consider merge join if tables have been sorted
 - Consider index join if index exists
- To pick the best, DBMS needs to maintain data statistics




Query Optimization

- $R(A, B) \quad S(B, C) \quad T(C, D)$:

SELECT * FROM R, S, T

WHERE $R.B = S.B$ AND $S.C = T.C$ AND $R.A = 10$ and $T.D < 30$

- Q: How can we process the above query?



Query Optimization

- Q: Focusing on just $R \bowtie S \bowtie T$, how many different ways?

- In general, for n way joins, $\frac{(2(n-1))!}{(n-1)!}$ ways
 - For $n = 10$, $\frac{18!}{9!} = 17 \times 10^9$ different ways!!!

Query Optimization

- In reality, picking the very best is too difficult
- DBMS tries to avoid “obvious mistakes” using a number of heuristics to examine only those plans that are likely to be good
 - Put the smallest table on the left
 - “Left-deep” tree
 - Push selection as deep as possible
 - ...
- For 90% of queries, DBMS picks a good query execution plan
 - To optimize the remaining 10%, companies pay big money to database consultants

Looking at Query Plan

- Many systems allow users to look at query plan
 - No SQL standard
 - Different systems use different syntax
- Examples
 - My SQL, PostgreSQL: EXPLAIN SELECT ...
 - Oracle: EXPLAIN PLAN FOR SELECT ...
 - MS SQL Server: SET SHOWPLAN_TEXT ON

Statistics Collection for DBMS

- “Cost-based optimizer”:
 - DBMS uses statistics on tables/indexes to pick the best query execution plan
 - Keeping correct stats is *very important.* Without correct stats, DBMS may do stupid things
- Oracle
 - ANALYZE TABLE <table> COMPUTE STATISTICS
 - ANALYZE TABLE <table> ESTIMATE STATISTICS ---- cheaper than COMPUTE
- DB2
 - RUN ON TABLE <userid>.<table> AND INDEXES ALL
- MySQL does not have a cost-based optimizer
 - Rule-based optimizer: Use simple heuristics only without looking at the actual data