

CS143: MongoDB (NoSQL)

Book Chapters

(7th) Chapter 10.2

MongoDB

- Database for JSON objects
 - “NoSQL database”
- Schema-less: no predefined schema
 - MongoDB will store anything with no complaint!
 - No normalization or joins
 - Use [Mongoose](#) for ensuring structure in the data
- Adopts JavaScript philosophy
 - “Laissez faire” policy
 - * Don’t be too strict! Handle user request in a “reasonable” way
 - Both blessing and curse

Document in MongoDB

- Data is stored as a *collection* of *documents*
 - *Document*: (almost) JSON object
 - *Collection*: group of “similar” documents
- Example

```
{
  "_id": ObjectId(8df38ad8902c),
  "title": "MongoDB",
  "description": "MongoDB is NoSQL database",
  "tags": ["mongodb", "database", "NoSQL"],
  "likes": 100,
  "comments": [
    { "user": "lover", "comment": "Great book!" },
    { "user": "hater", "comment": "Worst ever!" }
  ]
}
```

- `_id` field: primary key
 - Its value must be unique in the collection
 - May be of any type other than array
 - If not provided, `_id` is automatically added with a unique `ObjectId` value
- Stored as BSON (Binary representation of JSON)
 - Supports more data types than JSON
 - Does not require double quotes for field names
- Analogy
 - Document in MongoDB \approx row in RDB
 - Collection in MongoDB \approx table in RDB

MongoDB vs RDB

MongoDB document

- Preserves structure
 - Nested objects
- Potential redundancy
- Hierarchical view of a particular app
- Retrieving data with different “view” is difficult

RDB relation

- “Flattens” data
 - Set of flat rows
- Removes redundancy
- Flat schema based on the intrinsic nature of data
- Easy to obtain different “view” using efficient “joins”

Basic MongoDB Commands

- Basic administration
 - `mongo`: start MongoDB shell
 - `use <dbName>`: use the database
 - `show dbs`: show list of databases
 - `show collections`: show list of collections
 - `db.colName.drop()`: delete `colName` collection
 - `db.dropDatabase()`: delete current database
- CRUD operations
 - Create: `insertOne()`, `insertMany()`

- Retrieve: `findOne()`, `find()`
- Update: `updateOne()`, `updateMany()`
- Delete: `deleteOne()`, `deleteMany()`

MongoDB commands for CRUD

- Create: `insertX(doc(s))`

```
db.books.insertOne({title: "MongoDB", likes: 100})
db.books.insertMany([{title: "a"}, {title: "b"}])
```

- Retrieve: `findX(condition)`

```
db.books.findOne({likes: 100})
db.books.find({$and: [{likes: {$gte: 10}}, {likes: {$lt: 20}}]})
```

- `findOne()` returns the first (?) matching document for multiple matches
- Other boolean/comparison operators: `$or`, `$not`, `$gt`, `$ne`, ...

- Update: `updateX(condition, update_op)`

```
db.books.updateOne({title: "MongoDB"}, {$set: {title: "MongoDB II"}})
db.books.updateMany({title: "MongoDB"}, {$inc: {likes: 1}})
```

- Other update operators: `$mul` (multiply), `$unset` (remove the field), ...

- Delete: `deleteX(condition)`

```
db.books.deleteOne({title: "MongoDB"})
db.books.deleteMany({likes: {$lt: 100}})
```

MongoDB Queries: Aggregates

- MongoDB allows posing complex queries using “aggregates”
 - MongoDB aggregates \approx SQL select queries
 - An “aggregate pipeline” consists of multiple “aggregate stages”
 - * pipeline \approx select statement
 - * stage \approx select clause
- Example

```
{ _id: 1, cust_id: "a", status: "A", amount: 50 }
{ _id: 2, cust_id: "a", status: "A", amount: 100 }
{ _id: 3, cust_id: "c", status: "D", amount: 25 }
{ _id: 4, cust_id: "d", status: "C", amount: 125 }
{ _id: 5, cust_id: "d", status: "A", amount: 25 }
```

```
db.orders.aggregate([
  { $match: { status: "A" } },
  { $group: { _id: "$cust_id", total: { $sum: "$amount" }, count: {
    $sum: 1 } }},
  { $sort: { total: -1 } }
]);
```

- `$match` \approx where
 - `$group` \approx group by
 - * `_id` is the group by attribute
 - `$sort` \approx order by
 - `$limit` \approx fetch first
 - `$project` \approx select
 - `$unwind`: replicate document per every element in the array
 - * { `$unwind: "y"` } converts {`"x": 1, "y": [1, 2]` } to {`"x": 1, "y": 1`}, {`"x": 1, "y": 2` }
 - `$lookup`: “look up and join” another document based on attribute value
 - * {`$lookup: { from: <collection to join>, localField: <local join attr>, foreignField: <remote join attr>, as: <output field name> }`}
 - * matching documents are returned as an array in `<output field name>`
- More on MongoDB aggregates
 - Short tutorial: <https://studio3t.com/knowledge-base/articles/mongodb-aggregation-framework/>
 - Reference: <https://docs.mongodb.com/manual/reference/method/db.collection.aggregate/>

Index

- Indexes can be built for efficient retrieval
- `db.books.createIndex({title:1, likes:-1})`
 - Create one index on combined attributes “title” and “likes”
 - 1 means ascending order, -1 means descending order

More on MongoDB

- We learned just the basic
 - Enough for our project
- But MongoDB has many more features:
 - Aggregate queries
 - Transactions
 - Replication
 - (Auto)sharding

— ...

- Read MongoDB documentation and online tutorials to learn more