

CS 161 Fundamentals of Artificial Intelligence

Lecture 5

Local search algorithms

Quanquan Gu

Department of Computer Science
UCLA


January 20, 2022

Outline

- Hill-climbing
- Simulated annealing
- Local beam search
- Genetic algorithms
- Local search in continuous spaces


Iterative improvement algorithms

- In many optimization problems, **path** is irrelevant; the goal state itself is the solution




Iterative improvement algorithms

- ▶ In many optimization problems, **path** is irrelevant; the goal state itself is the solution
- ▶ For example, in the 8-queens problem, what matters is the final configuration of queens, not the order in which they are added.



Iterative improvement algorithms


- ▶ In many optimization problems, **path** is irrelevant; the goal state itself is the solution
- ▶ For example, in the 8-queens problem, what matters is the final configuration of queens, not the order in which they are added.



- ▶ In such cases, can use **iterative improvement** algorithms; keep a single “current” state, try to improve it

Iterative improvement algorithms


- ▶ In many optimization problems, **path** is irrelevant; the goal state itself is the solution
- ▶ For example, in the 8-queens problem, what matters is the final configuration of queens, not the order in which they are added.



- ▶ In such cases, can use **iterative improvement** algorithms; keep a single “current” state, try to improve it
- ▶ Constant space, suitable for online as well as offline search

Example: Travelling Salesperson Problem


- ▶ Goal: to find the shortest path that visits each city and returns to the origin city
- ▶ Start with any complete tour (may have cross path, not optimal)
- ▶ Perform pairwise exchanges, each iteration reduces length of path



Variants of this approach get within 1% of optimal very quickly with thousands of cities


Example: n -queens

- ▶ Goal: Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal
- ▶ Move a queen to reduce number of conflicts



Almost always solves n -queens problems almost instantaneously for very large n , e.g., $n = 1\text{million}$

State space landscape



- ▶ Goal: to find global maximum
- ▶ Complete: finds a goal if one exists;
- ▶ Optimal: finds a global minimum/maximum
- ▶

Hill-climbing (or gradient ascent/descent)

“Like climbing Everest in thick fog with amnesia”


- Moves in the direction of increasing value

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
  end
```


Hill-climbing contd.

Useful to consider **state space landscape**



- ▶ Escape from shoulders: **Random sideways moves**, maybe loop on flat maxima
- ▶ Escape from local maxima: **Random-restart hill climbing**, trivially complete

Hill-climbing example¹




$h(x) = +1$ for all the blocks in the support structure if the block is correctly positioned otherwise -1 for all the blocks in the support structure.

$h(1)$: A is incorrectly positioned with 3 support blocks (-3), B is incorrectly positioned with 0 support blocks (-0), C is incorrectly positioned with 1 support blocks (-1), D is incorrectly positioned with 2 support blocks (-2).

¹Reference: <https://www.baeldung.com/java-hill-climbing-algorithm>


Hill-climbing example



$h(x) = +1$ for all the blocks in the support structure if the block is correctly positioned otherwise -1 for all the blocks in the support structure.

- ▶ $h(1) = (-3) + (-0) + (-1) + (-2) = -6$
- ▶ $h(2) = (+0) + (-0) + (-1) + (-2) = -3$
- ▶ $h(3) = (+0) + (-0) + (-1) + (-0) = -1$
- ▶ ...

Hill-climbing example: 8 queen



- ▶ Each state has 8 queens on the board, one per column
- ▶ Successors: move a single queen to another square in the same column
- ▶ heuristic cost function h : the number of pairs of queens that are attacking each other
- ▶ Before: $h = 0(\text{column}) + 5(\text{row}) + 12(\text{diagonal}) = 17$
- ▶ After: $h = 0(\text{column}) + 4(\text{row}) + 8(\text{diagonal}) = 12$

Simulated annealing

Idea: escape local maxima by allowing some “bad” moves
but gradually decrease their size and frequency

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
           schedule, a mapping from time to “temperature”
  local variables: current, a node
                   next, a node
                   T, a “temperature” controlling prob. of downward steps

  current ← MAKE-NODE(INITIAL-STATE[problem])
  for t ← 1 to ∞ do
    T ← schedule[t]
    if T = 0 then return current
    next ← a randomly selected successor of current
     $\Delta E \leftarrow \text{VALUE}[\textit{next}] - \text{VALUE}[\textit{current}]$ 
    if  $\Delta E > 0$  then current ← next
    else current ← next only with probability  $e^{\Delta E/T}$ 
```

Properties of simulated annealing

- ▶ If the move improves the situation, it is always accepted. Otherwise, the algorithm accepts the move with some probability less than 1.
- ▶ The probability decreases exponentially with the “badness” of the move—the amount ΔE by which the evaluation is worsened.
- ▶ The probability also decreases as the “temperature” T goes down: “bad” moves are more likely to be allowed at the start when T is high, and they become more unlikely as T decreases.
- ▶ If the schedule lowers T slowly enough, the algorithm will find a global optimum with probability approaching 1

Properties of simulated annealing

At fixed “temperature” T , state occupation probability reaches Boltzman distribution

$$p(x) = \alpha e^{-\frac{E(x)}{kT}}$$

T decreased slowly enough \implies always reach best state x^*
($E(x^*) = \max_x E(x)$)

- $e^{-\frac{E(x^*)}{kT}} / e^{-\frac{E(x)}{kT}} = e^{\frac{E(x^*) - E(x)}{kT}} \gg 1$ for small T
- Thus, very likely to choose x^* !

Local beam search

- **Idea**: keep k states instead of 1; At each step, all the successors of all k states are generated. If any one is a goal, the algorithm halts. Otherwise, it selects the best k successors from the complete list and repeats.

- Not the same as k searches run in parallel!

Searches that find good states recruit other searches to join them


- **Problem**: quite often, all k states end up on same local hill

- **Stochastic beam search**: choose k successors randomly, biased towards good ones

Observe the close analogy to natural selection!

Evolutionary algorithms/Genetic algorithms

- stochastic local beam search + generate successors from **pairs** of states




- Fitness function: higher score, higher chance to be selected
- Cross-over: crossover point is chosen randomly
- Mutation: small probability

Genetic algorithms contd.

GAs require states encoded as strings

Crossover helps **iff substrings are meaningful components**



Local search in continuous state spaces

Suppose we want to site three airports in Romania:

- 6-D state space defined by $(x_1, y_1), (x_2, y_2), (x_3, y_3)$
- objective function $f(x_1, y_1, x_2, y_2, x_3, y_3) =$
sum of squared distances from each city to nearest airport

Discretization methods turn continuous space into discrete space,
e.g., **empirical gradient** considers $\pm\delta$ change in each coordinate

Gradient methods

Gradient methods compute

$$\nabla f = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial y_1}, \frac{\partial f}{\partial x_2}, \frac{\partial f}{\partial y_2}, \frac{\partial f}{\partial x_3}, \frac{\partial f}{\partial y_3} \right)$$

to increase/reduce f , e.g., by $\mathbf{x} \leftarrow \mathbf{x} + \alpha \nabla f(\mathbf{x})$

Newton-Raphson Method

Sometimes can solve for $\nabla f(\mathbf{x}) = 0$ exactly.

Newton-Raphson (1664, 1690) iterates $\mathbf{x} \leftarrow \mathbf{x} - \mathbf{H}_f^{-1}(\mathbf{x}) \nabla f(\mathbf{x})$
to solve $\nabla f(\mathbf{x}) = 0$, where $\mathbf{H}_{ij} = \partial^2 f / \partial x_i \partial x_j$

Acknowledgment

The slides are adapted from Stuart Russell et al.