

Lecture 2:

Overview Fall 2022

Kai-Wei Chang
CS @ UCLA

kw+cm146@kwchang.net

The instructor gratefully acknowledges Eric Eaton (UPenn), who assembled the original slides, Jessica Wu (Harvey Mudd), David Kauchak (Pomona), Dan Roth (Upenn), Sriram Sankararaman (UCLA), whose slides are also heavily used, and the many others who made their course materials freely available online.

Announcement

- ❖ There is a discussion session on Friday
 - ❖ See session/time/loc at myUCLA
- ❖ Math Review Quiz is on BruinLearn

This Lecture

- ❖ Learning Protocols
 - ❖ Supervised Learning
 - ❖ Unsupervised Learning
- ❖ Challenges in ML
- ❖ Framing Learning Problems

Type of learning protocols

Supervised Learning

Training phase:



lion




Not lion

Test phase:




??

Supervised Learning



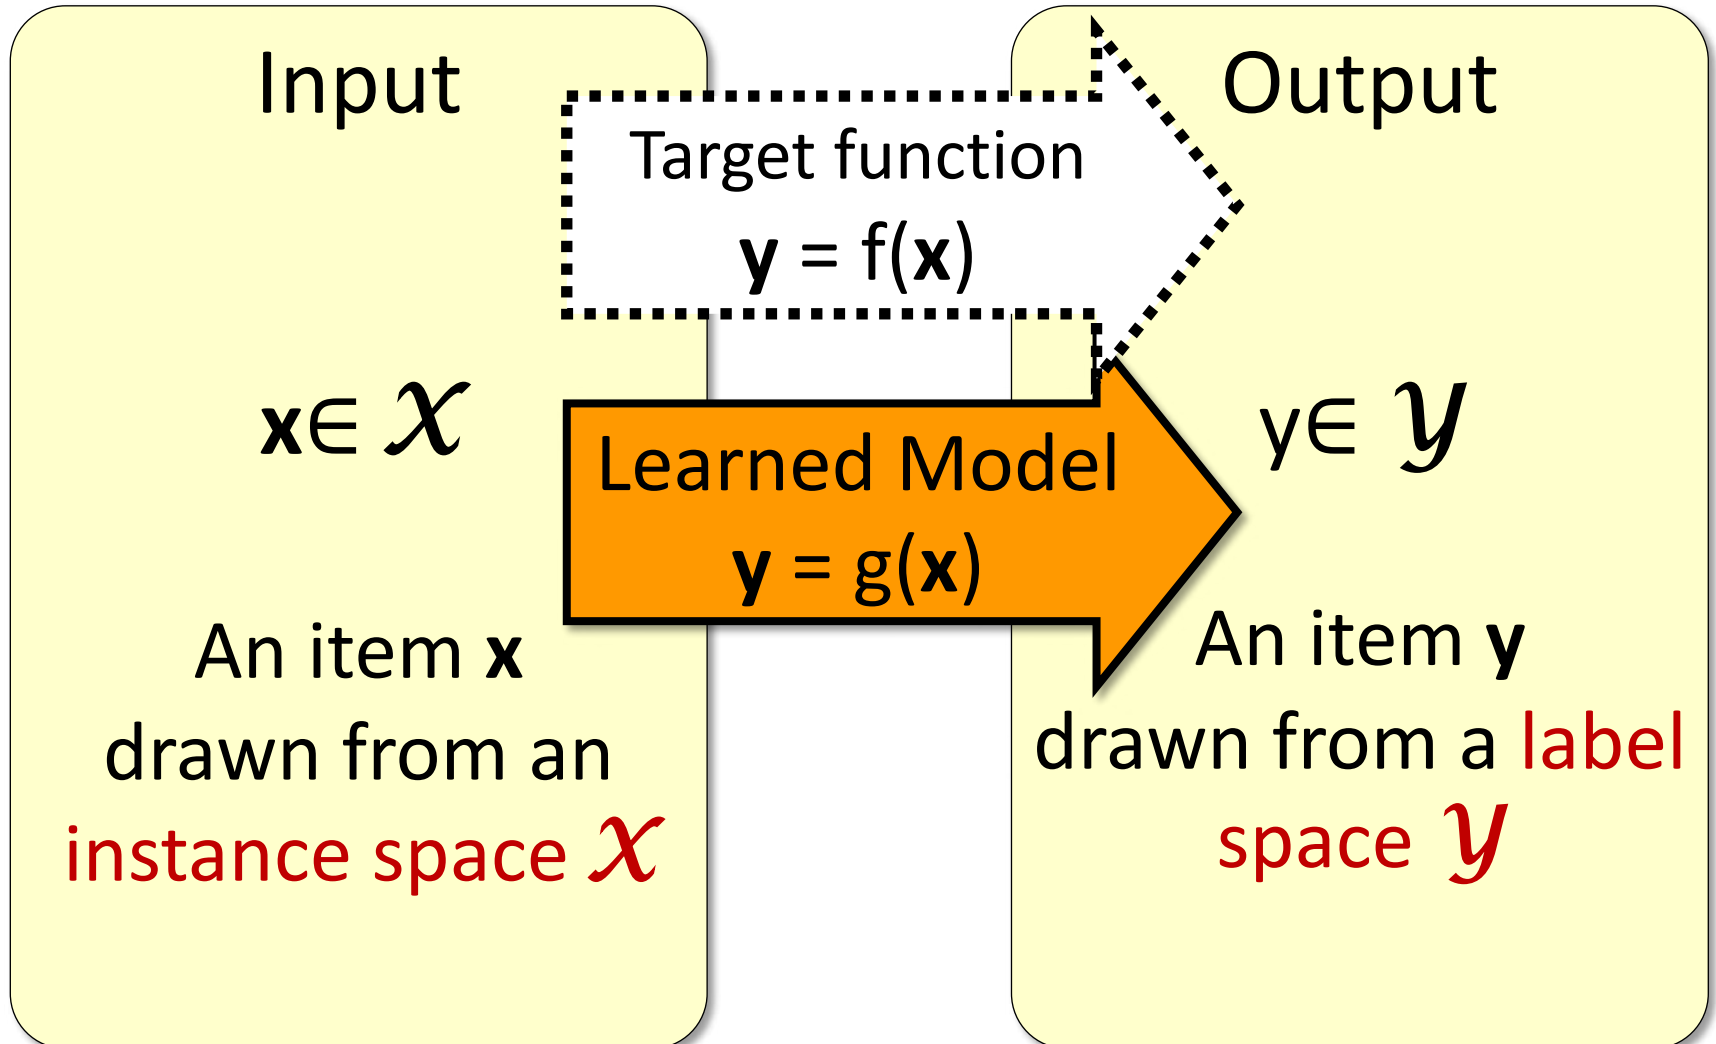
- ❖ We consider systems that apply a function $f()$ to input items \mathbf{x} and return an output $\mathbf{y} = f(\mathbf{x})$.

Supervised Learning

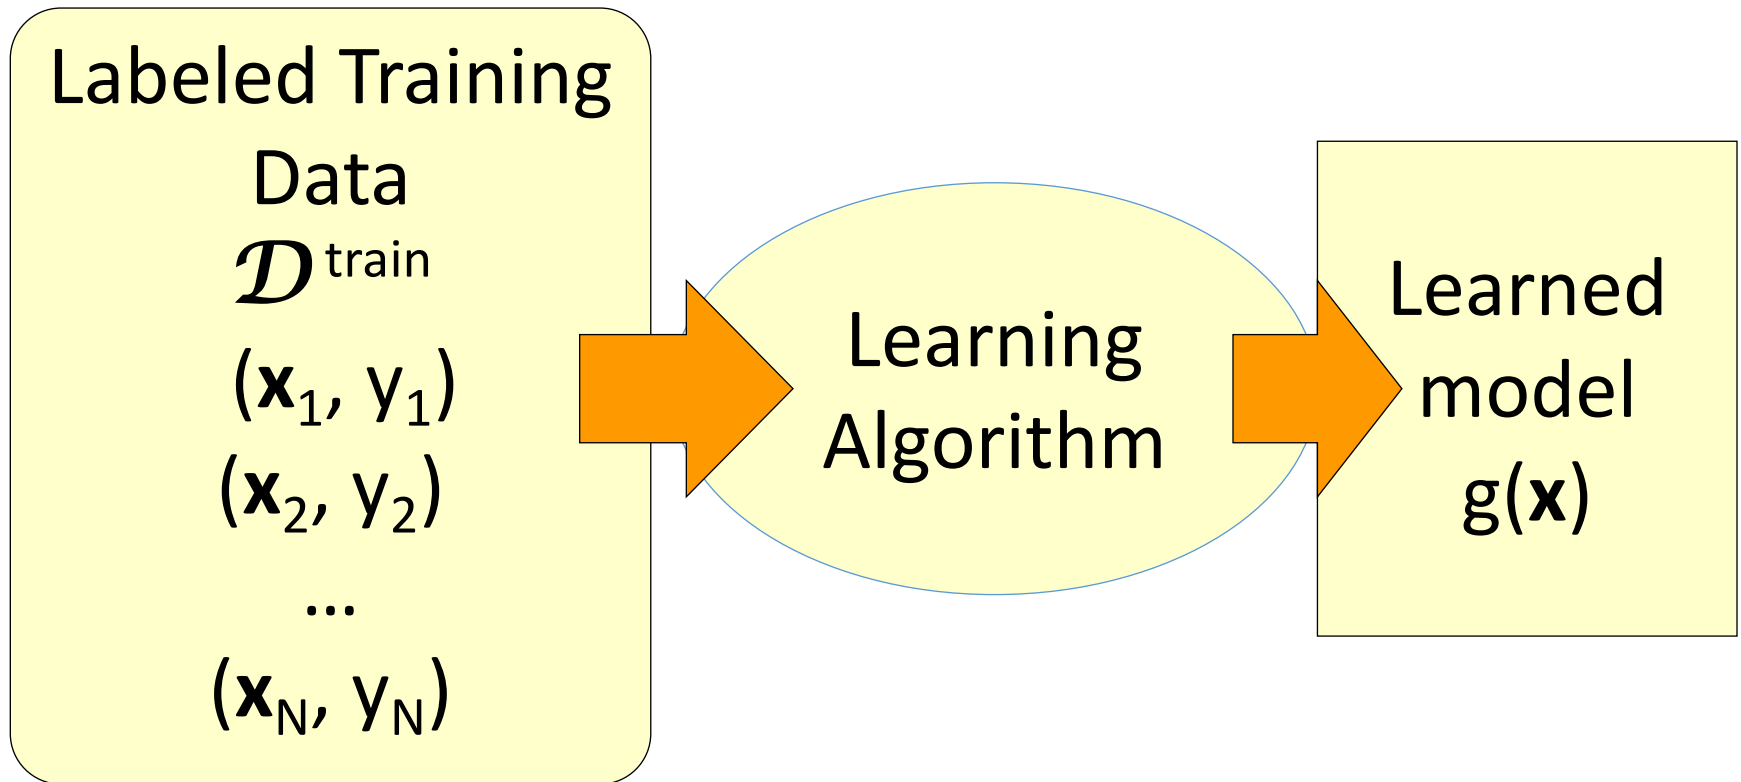


- ❖ In (supervised) machine learning, we deal with systems whose $f(\mathbf{x})$ is learned from examples.

Supervised Learning



Supervised Learning: Training



- ❖ Give the learner examples in $\mathcal{D}^{\text{train}}$
- ❖ The learner returns a model $g(\mathbf{x})$

Supervised Learning: Testing

Labeled
Test Data

$\mathcal{D}^{\text{test}}$

(\mathbf{x}'_1, y'_1)

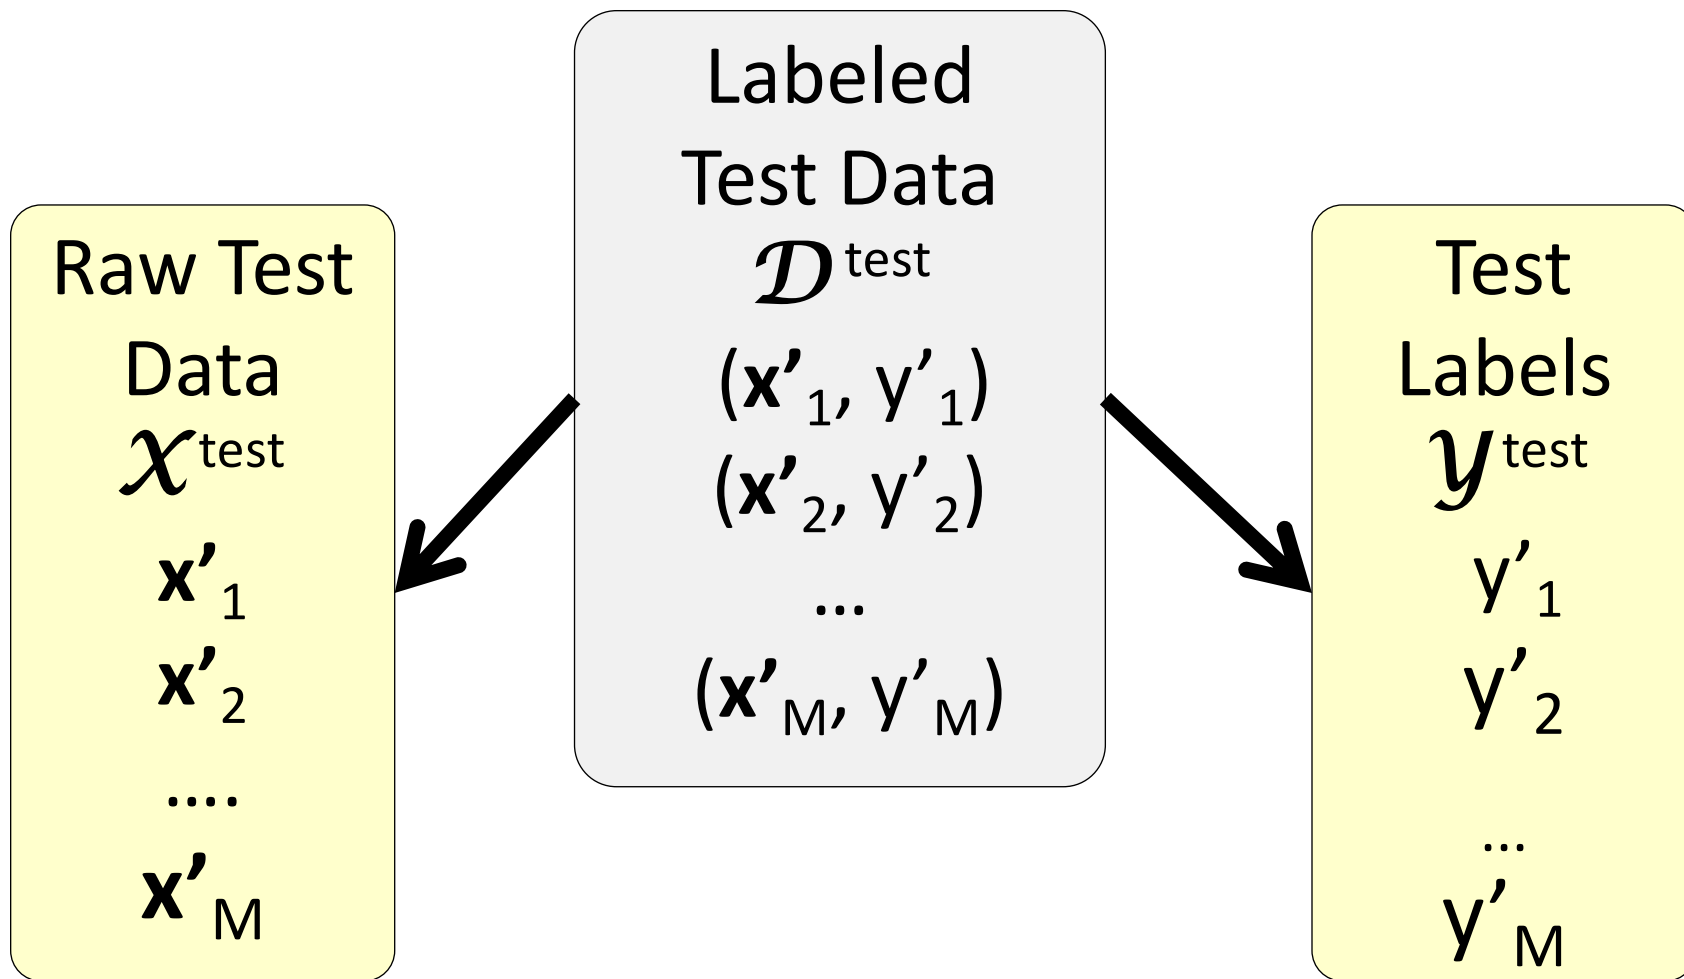
(\mathbf{x}'_2, y'_2)

...

(\mathbf{x}'_M, y'_M)

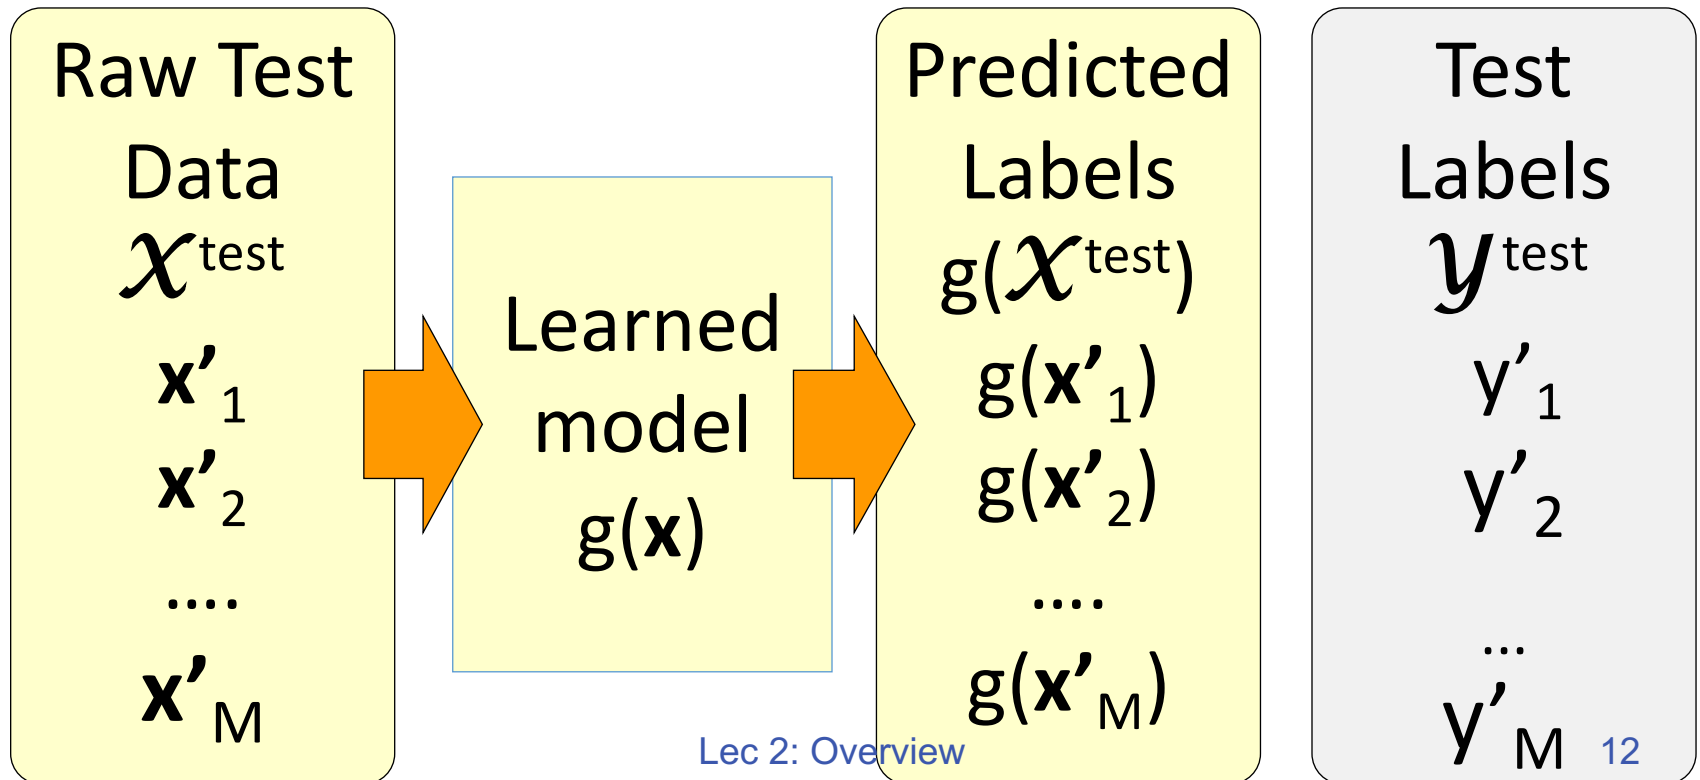
- ❖ Reserve some labeled data for testing

Supervised Learning: Testing




Supervised Learning: Testing

- ❖ Apply the model to the raw test data
- ❖ Evaluate by comparing predicted labels against the test labels




Unsupervised learning


- ❖ Given: **unlabeled** inputs
- ❖ Goal: learn some intrinsic structure in inputs




How many “kinds of monsters” are there?




How many “kinds of monsters” are there?



How many “kinds of monsters” are there?



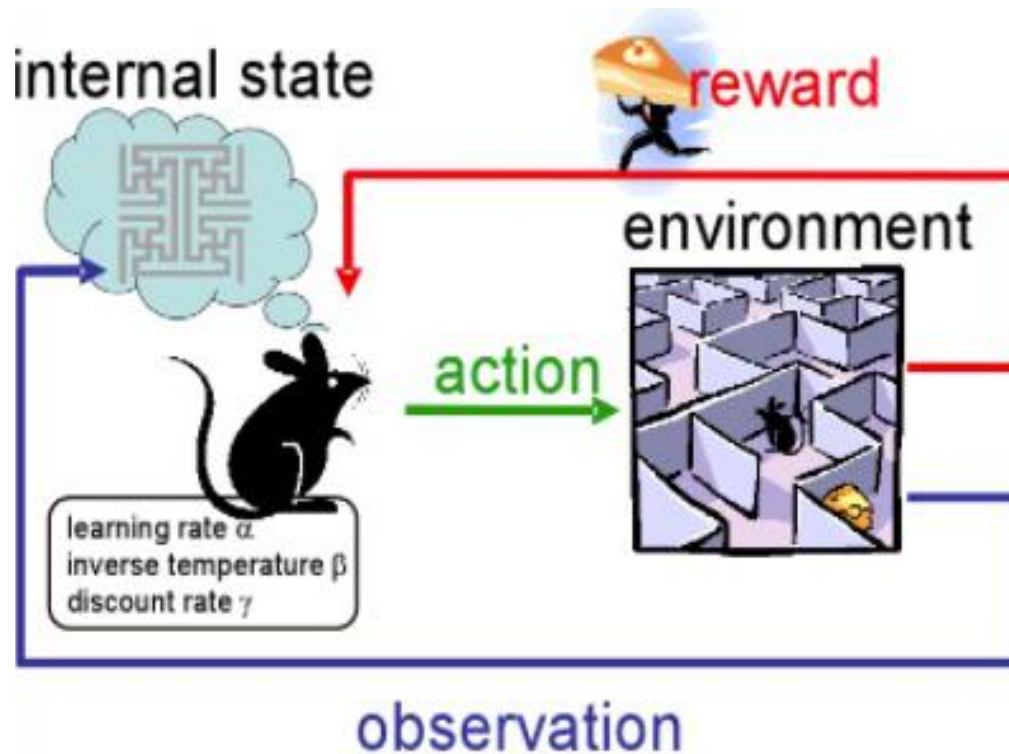
Decipher



Credit: Dan Roth



Reinforcement Learning

- ❖ Given sequences of states and actions with rewards
- ❖ Learn policy that maximizes agent's reward




Challenges in ML

Structured Inference

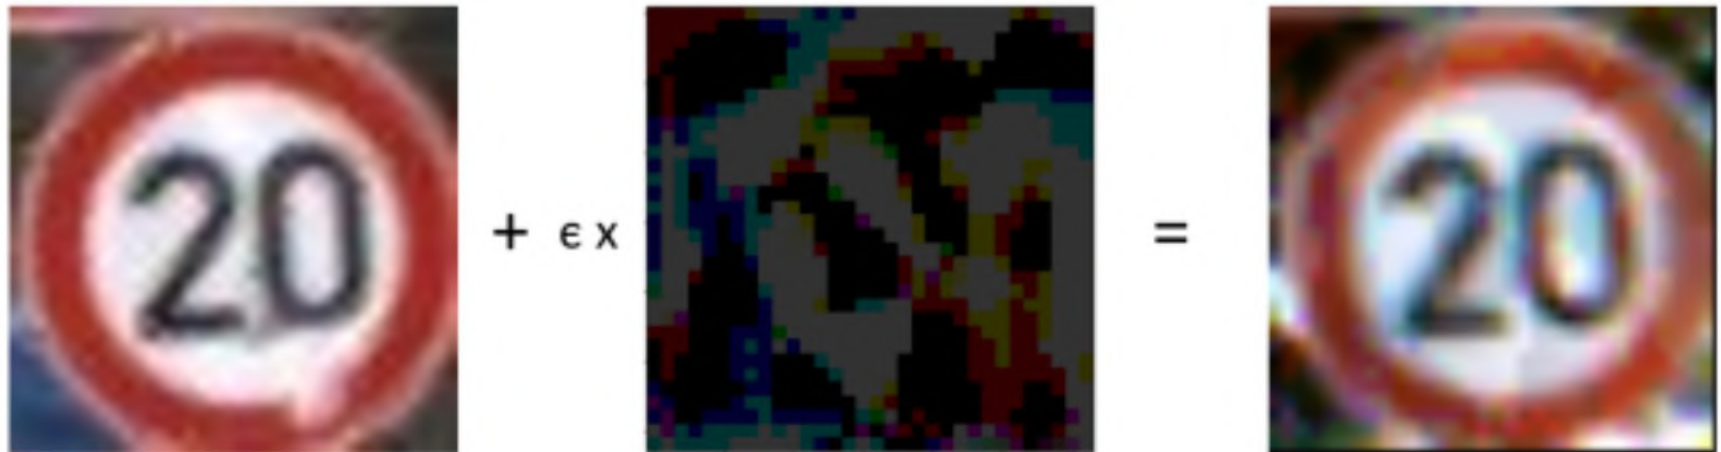


Robustness

Car or shoe?



Adversarial Attack



93%, 20 Km/h Sign + ϵx = 90%, 80 Km/h Sign

$sign(\nabla * J(\theta, x, y))$



<https://arxiv.org/abs/1712.09327v1>


Commonsense

❖ Winograd Schema (1972)

The city councilmen refused the demonstrators a permit because they feared violence.

The city councilmen refused the demonstrators a permit because they advocated violence.

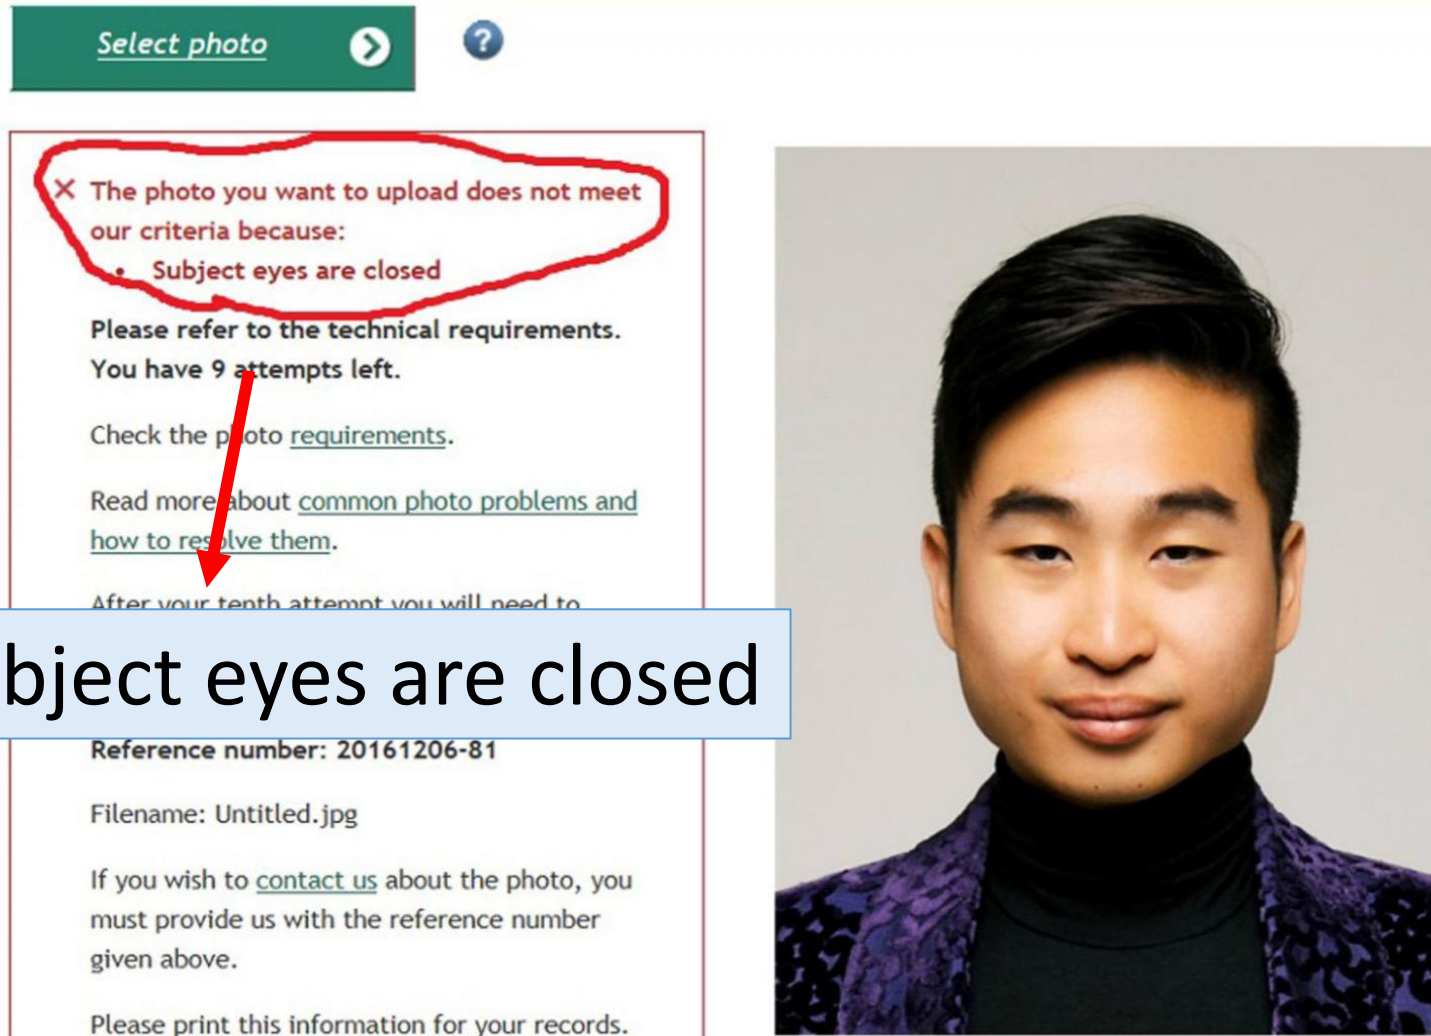
❖ Visual Commonsense



Is it raining outside?

- a) Yes, it is snowing.
- b) Yes, [person8] and [person10] are outside.
- c) No, it looks to be fall.
- d) Yes, it is raining heavily.

Fairness/Inclusion in ML

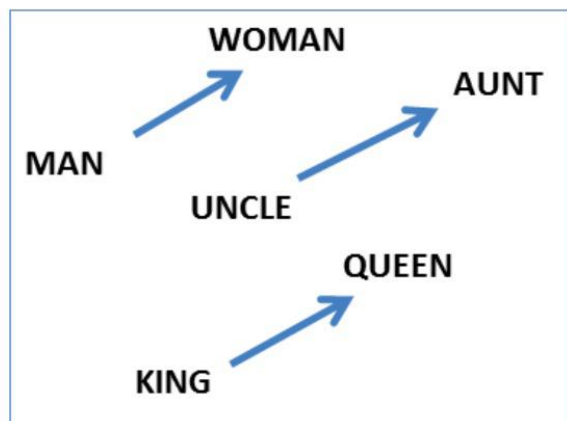


A screenshot of New Zealand man Richard Lee's passport photo rejection notice, supplied to Reuters December 7, 2016. Richard Lee/Handout via REUTERS

LEE: R. HOUTO

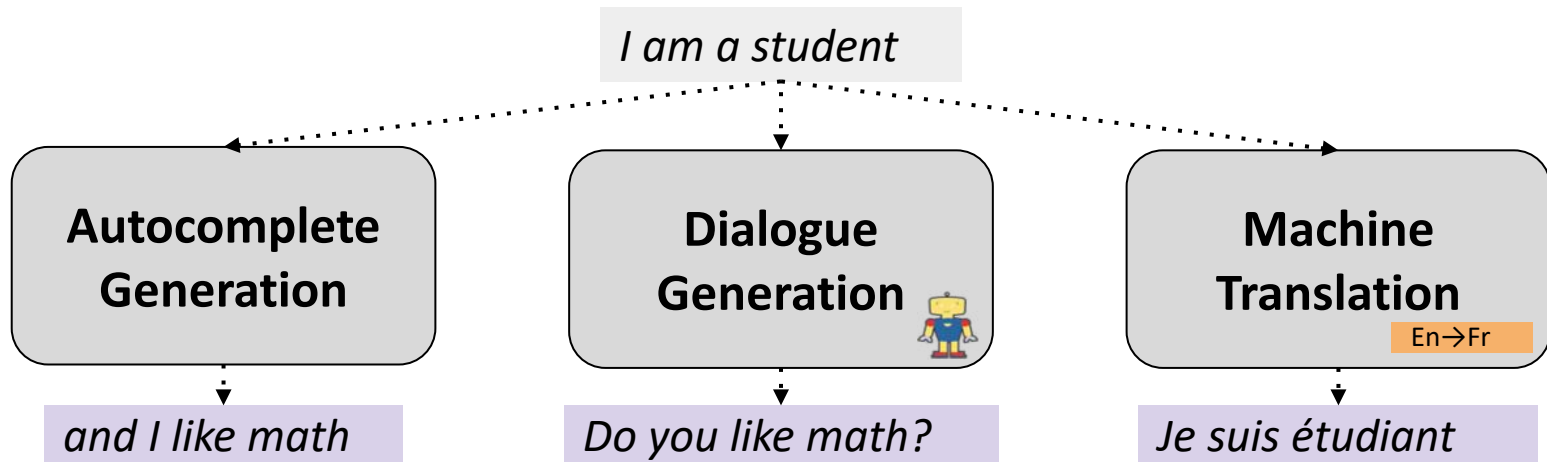
Fairness in ML-- Word embedding bias

$$\diamond v_{man} - v_{woman} + v_{uncle} \sim v_{aunt}$$

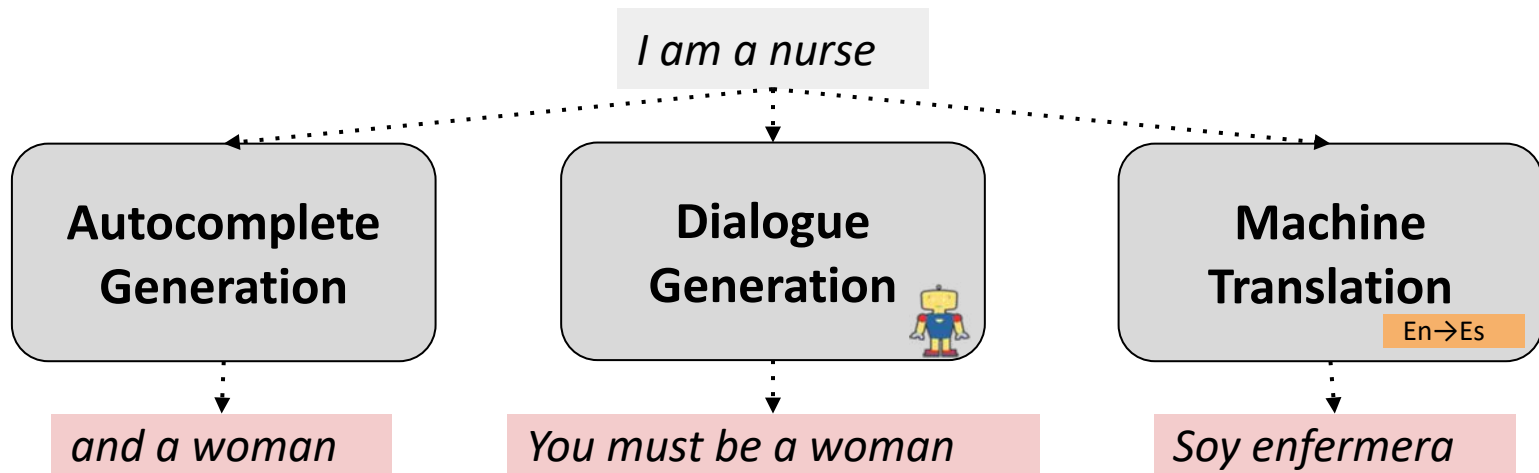


he: ____	she: ____
uncle	aunt
lion	
surgeon	
architect	
beer	
professor	

We use Google w2v embedding trained from the news



Language generations can be gendered!



Societal Biases in Language Generation: Progress and Challenges

Emily Sheng, Kai-Wei Chang, Prem Natarajan, and Nanyun Peng, in ACL, 2021.

Misgendering in NLG

Alex went to the hospital for their appointment. [MASK] felt sick.

Prediction	Score
Alex went to the hospital for their appointment . She felt sick .	45.3%
Alex went to the hospital for their appointment . He felt sick .	36%
Alex went to the hospital for their appointment . Alex felt sick .	5.8%
Alex went to the hospital for their appointment . I felt sick .	2.3%
Alex went to the hospital for their appointment . They felt sick .	0.4%

<https://demo.allennlp.org/masked-lm>

Harms of Gender Exclusivity and Challenges in Non-Binary Representation in Language Technologies

Sunipa Dev, Masoud Monajatipoor, Anaelia Ovalle, Arjun Subramonian, Jeff Phillips, and Kai-Wei Chang, in EMNLP, 2021.

Framing a Learning Problem

How we set up a learning problem

❖ The Badges Game

- ❖ This is an example of the key learning protocol:
supervised learning

The Badges game

+ Naoki Abe

- Eric Baum

- ❖ Conference attendees to the 1994 Machine Learning conference were given **name badges labeled with + or -**.
- ❖ What function was used to assign these labels?

Training data

- | | | |
|---------------------|-------------------|--------------------|
| + Naoki Abe | + Peter Bartlett | + Carla E. Brodley |
| - Myriam Abramson | - Eric Baum | + Nader Bshouty |
| + David W. Aha | + Welton Becket | - Wray Buntine |
| + Kamal M. Ali | - Shai Ben-David | - Andrey Burago |
| - Eric Allender | + George Berg | + Tom Bylander |
| + Dana Angluin | + Neil Berkman | + Bill Byrne |
| - Chidanand Apte | + Malini Bhandaru | - Claire Cardie |
| + Minoru Asada | + Bir Bhanu | + John Case |
| + Lars Asker | + Reinhard Blasig | + Jason Catlett |
| + Javed Aslam | - Avrim Blum | - Philip Chan |
| + Jose L. Balcazar | - Anselm Blumer | - Zhixiang Chen |
| - Cristina Baroglio | + Justin Boyan | - Chris Darken |

Raw test data

Gerald F. DeJong
Chris Drummond
Yolanda Gil
Attilio Giordana
Jiarong Hong
J. R. Quinlan

Priscilla Rasmussen
Dan Roth
Yoram Singer
Lyle H. Ungar

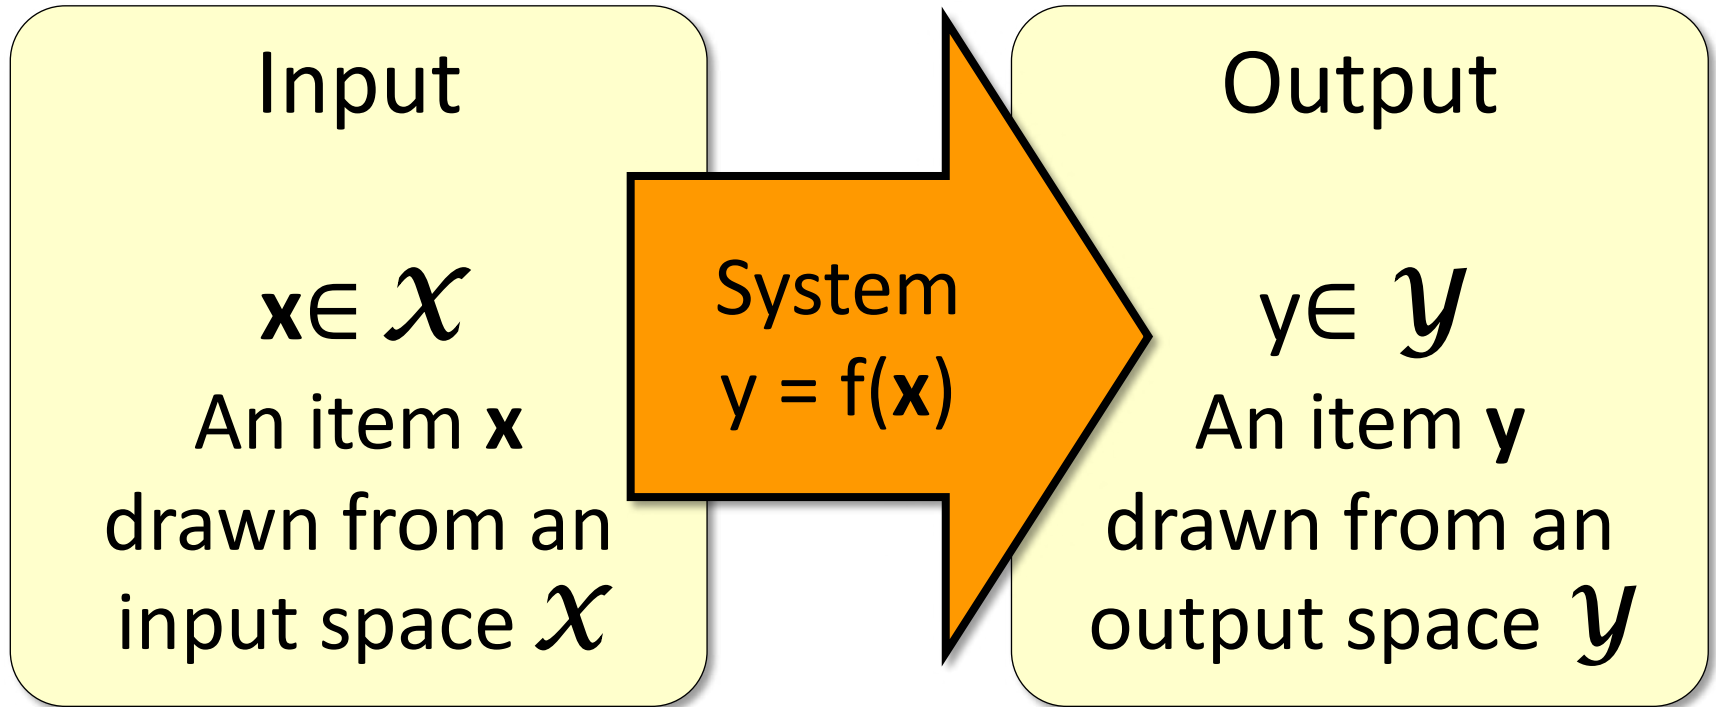
Labeled test data

- + Gerald F. DeJong
- Chris Drummond
- + Yolanda Gil
- Attilio Giordana
- + Jiarong Hong
- J. R. Quinlan
- Priscilla Rasmussen
- + Dan Roth
- + Yoram Singer
- Lyle H. Ungar

Exercise: What is the rule?

- | | | |
|---------------------|-------------------|--------------------|
| + Naoki Abe | + Peter Bartlett | + Carla E. Brodley |
| - Myriam Abramson | - Eric Baum | + Nader Bshouty |
| + David W. Aha | + Welton Becket | - Wray Buntine |
| + Kamal M. Ali | - Shai Ben-David | - Andrey Burago |
| - Eric Allender | + George Berg | + Tom Bylander |
| + Dana Angluin | + Neil Berkman | + Bill Byrne |
| - Chidanand Apte | + Malini Bhandaru | - Claire Cardie |
| + Minoru Asada | + Bir Bhanu | + John Case |
| + Lars Asker | + Reinhard Blasig | + Jason Catlett |
| + Javed Aslam | - Avrim Blum | - Philip Chan |
| + Jose L. Balcazar | - Anselm Blumer | - Zhixiang Chen |
| - Cristina Baroglio | + Justin Boyan | - Chris Darken |

Supervised Learning



- ❖ We consider systems that apply a function $f()$ to input items \mathbf{x} and return an output $y = f(\mathbf{x})$.

Using supervised learning

- ❖ What is our **instance space**?
 - ❖ Gloss: What kind of features are we using?
- ❖ What is our **label space**?
 - ❖ Gloss: What kind of learning task are we dealing with?
- ❖ What is our **hypothesis space**?
 - ❖ Gloss: What kind of functions (models) are we learning?
- ❖ What **learning algorithm** do we use?
 - ❖ Gloss: How do we learn the model from the labeled data?
- ❖ What is our **loss function**/evaluation metric?
 - ❖ Gloss: How do we measure success? What drives learning?

1. Input: The instance space \mathcal{X}

Input

$\mathbf{x} \in \mathcal{X}$

An item \mathbf{x}
drawn from an
instance space
 \mathcal{X}

\mathbf{x} is represented in a feature space

- Typically $\mathbf{x} \in \{0,1\}^n$ or \mathbb{R}^N
- Usually represented as a vector
- We call it input vector

Example:

Boolean features:

Does this email contain the word 'money'?

Numerical features:

How often does 'money' occur in this email

What is the width/height of this bounding box?

What is the length of the first name?

What's \mathcal{X} for the Badges game?

❖ Possible features:

- Length of their first or last name?
- Does the name contain letter 'x'?
- How many vowels does their name contain?
- Is the n-th letter a vowel?

+ Naoki Abe

- Myriam Abramson

+ David W. Aha

+ Kamal M. Ali

- Eric Allender

+ Dana Angluin

+ Peter Bartlett

- Eric Baum

+ Welton Becket

- Shai Ben-David

+ George Berg

+ Neil Berkman

+ Carla E. Brodley

+ Nader Bshouty

- Wray Buntine


- Andrey Burago

+ Tom Bylander

+ Bill Byrne

\mathcal{X} as a vector space

- ❖ \mathcal{X} is an N-dimensional vector space (e.g. \mathbb{R}^N)
 - ❖ Each dimension = one feature.
- ❖ Each \mathbf{x} is a **feature vector** (hence the boldface \mathbf{x}).
- ❖ Think of $\mathbf{x} = [x_1 \dots x_N]$ as a point in \mathcal{X} :



Example: the badge game

+ Naoki Abe	+ Peter Bartlett	+ Carla E. <u>Brodley</u>
- <u>Myriam Abramson</u>	- Eric Baum	+ Nader <u>Bshouty</u>
+ David W. Aha	+ <u>Welton Becket</u>	- Wray <u>Buntine</u>
+ Kamal M. Ali	- <u>Shai Ben-David</u>	- <u>Andrey Burago</u>
- Eric <u>Allender</u>	+ George Berg	+ Tom <u>Bylander</u>
+ Dana <u>Angluin</u>	+ Neil <u>Berkman</u>	+ Bill Byrne

[first-char is vowel, first-char is A, first-char is N, second-char is vowel ...]

+ Naoki Abe

[0 , 0 , 1 , 1 ...]

- Avrim Blum

[1 , 1 , 0 , 0 ...]

Good features are essential

- ❖ The choice of features is crucial for how well a task can be learned.
 - ❖ In many application areas (language, vision, etc.), a lot of work goes into designing suitable features.
 - ❖ This requires domain expertise.
- ❖ CM146 can't teach you what specific features to use for your task.
 - ❖ But we will touch on some general principles

2. Output space

y is represented in **output space**
(**label space**)

Different kinds of output:

- Binary classification:
 $y \in \{-1, 1\}$
- Multiclass classification:
 $y \in \{1, 2, 3, \dots, K\}$
- Regression:
 $y \in \mathbb{R}$
- Structured output
 $y \in \{1, 2, 3, \dots, K\}^N$

Output

$$y \in \mathcal{Y}$$

An item y
drawn from a **label**
space \mathcal{Y}

Supervised Learning : Examples

Animal recognition

- ❖ x : Bitmap picture of the animal
- ❖ y :



Lion? Yes/**No**

Binary output



Lion/Cat/**Dog**

Multiclass output

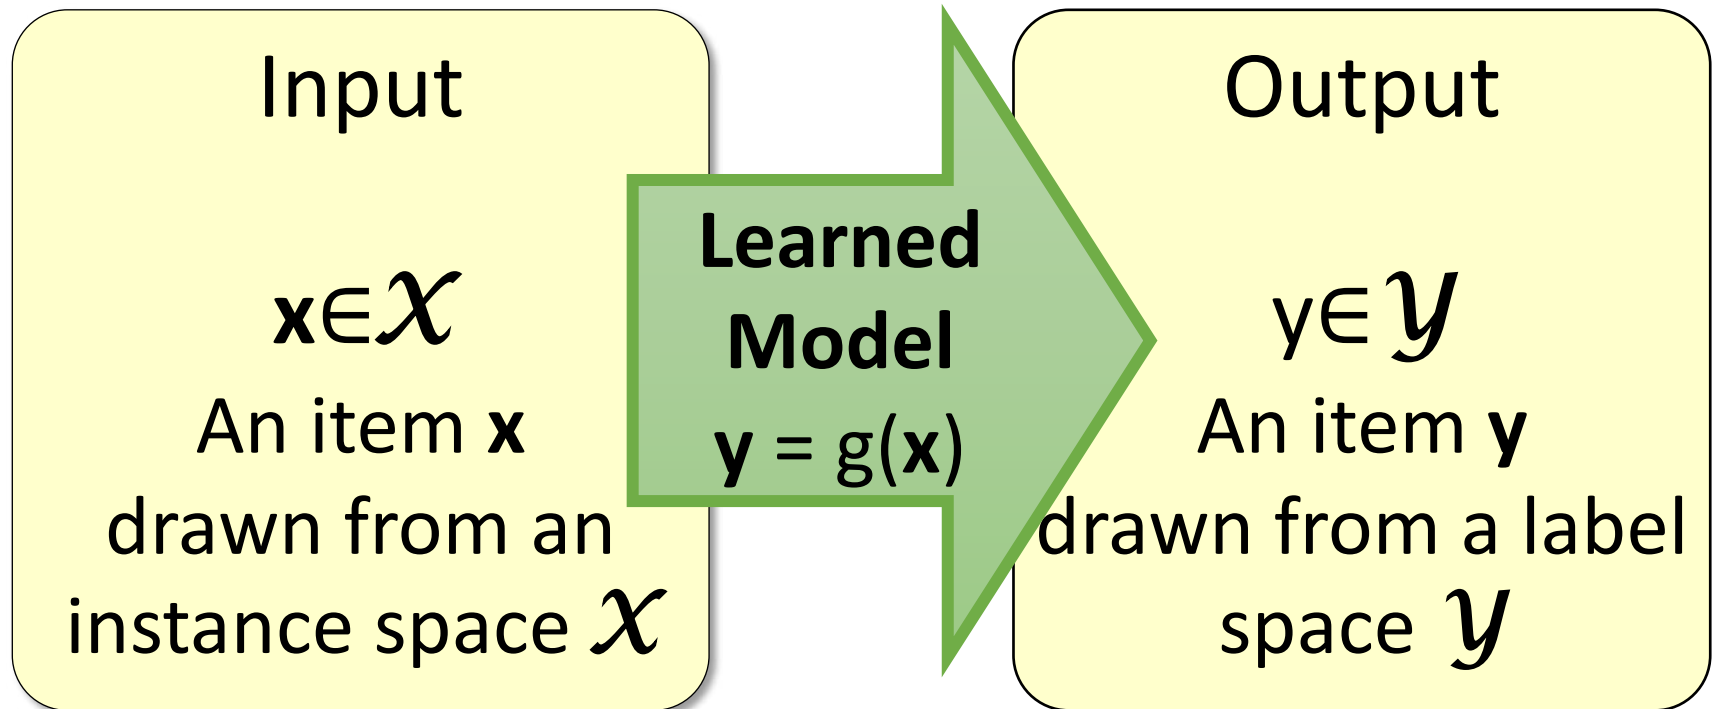


Lion/**Mammal**/**Dog**/Fish

Multilabel output


Output of the applications may differ from the output of ML models.

3. The model $g(\mathbf{x})$




- ❖ We need to choose what *kind* of model we want to learn

3. The model $g(\mathbf{x})$



- ❖ We need to choose what *kind* of model we want to learn

Boolean Function



x_1	x_2	y
0	0	0
0	1	0
1	0	1
1	1	1

Function 1

x_1	x_2	y
0	0	1
0	1	0
1	0	0
1	1	1

Function 2


x_1	x_2	y
0	0	0
0	1	0
1	0	0
1	1	0

Function 3

...

Hypothesis Space

A Learning Problem




Example	x_1	x_2	x_3	x_4	y
1	0	0	1	0	0
2	0	1	0	0	0
3	0	0	1	1	1
4	1	0	0	1	1
5	0	1	1	0	0
6	1	1	0	0	0
7	0	1	0	1	0

Can you learn this function?

What is it?

A function g is consistent to a dataset $D = \{(x_i, y_i)\}$ if $g(x_i) = y_i, \forall i$

Discussion: A Learning Problem



Example	x_1	x_2	x_3	x_4	y
1	0	0	1	0	0
2	0	1	0	0	0
3	0	0	1	1	1
4	1	0	0	1	1
5	0	1	1	0	0
6	1	1	0	0	0
7	0	1	0	1	0

Can you learn this function?

What is it?

A function g is consistent to a dataset $D = \{(x_i, y_i)\}$ if $g(x_i) = y_i, \forall i$

How many possible functions over four features?

How many function is consistent to D on the left

Hypothesis Space

How many possible functions over four features?

Complete Ignorance:

There are $2^{16} = 65536$ possible functions over four input features.

Example	X1	X2	X3	X4	y
	0	0	0	0	?
	0	0	0	1	?
	0	0	1	0	?
	0	0	1	1	?
	0	1	0	0	?
	0	1	0	1	?
	0	1	1	0	?
	0	1	1	1	?
	1	0	0	0	?
	1	0	0	1	?
	1	0	1	0	?
	1	0	1	1	?
	1	1	0	0	?
	1	1	0	1	?
	1	1	1	0	?
	1	1	1	1	?

Hypothesis Space

Complete Ignorance:

There are $2^{16} = 65536$ possible functions over four input features.

We can't figure out which one is correct until we've seen every possible input-output pair.

After observing seven examples we still have 2^9 possibilities for f

Is Learning Possible?

which one is the most likely one?

Example	X1	X2	X3	X4	y
	0	0	0	0	?
	0	0	0	1	?
	0	0	1	0	0
	0	0	1	1	1
	0	1	0	0	0
	0	1	0	1	0
	0	1	1	0	0
	0	1	1	1	?
	1	0	0	0	?
	1	0	0	1	1
	1	0	1	0	?
	1	0	1	1	?
	1	1	0	0	0
	1	1	0	1	?
	1	1	1	0	?
	1	1	1	1	?

Hypothesis Space

Complete Ignorance:

There are $2^{16} = 65536$ possible functions over four input features.

Example	X1	X2	X3	X4	y
	0	0	0	0	?
	0	0	0	1	?
	0	0	1	0	0
					1
					0
					0
					0
					?
					?
					1
					?
					?
					0
	1	1	0	0	?
	1	1	0	1	?
	1	1	1	0	?
	1	1	1	1	?

- There are $|Y|^{|\mathbf{X}|}$ possible functions $f(\mathbf{x})$ from the instance space \mathbf{X} to the label space \mathbf{Y} .
- Learners typically consider *only a subset* of the functions from \mathbf{X} to \mathbf{Y} , called the hypothesis space \mathbf{H} . $\mathbf{H} \subseteq |Y|^{|\mathbf{X}|}$

Is Learning Possible?

Hypothesis Space (2)

Simple Rules: **conjunctive rules**

of the form $y = x_i \wedge x_j \wedge \dots \wedge x_k$

e.g., $y = x_2 \wedge x_3$

$y = x_1 \wedge x_2 \wedge x_4$

How large is the hypothesis space?

Hypothesis Space (2)

1	0	0	1	0	0
2	0	1	0	0	0
3	0	0	1	1	1
4	1	0	0	1	1
5	0	1	1	0	0
6	1	1	0	0	0
7	0	1	0	1	0

Simple Rules: There are only 16 simple **conjunctive rules**

of the form $y = x_i \wedge x_j \wedge x_k$

<u>Rule</u>	<u>Counterexample</u>
-------------	-----------------------

$y = 1$

x_1

x_2

x_3

x_4

$x_1 \wedge x_2$

$x_1 \wedge x_3$

$x_1 \wedge x_4$

<u>Rule</u>	<u>Counterexample</u>
-------------	-----------------------

$x_2 \wedge x_3$

$x_2 \wedge x_4$

$x_3 \wedge x_4$

$x_1 \wedge x_2 \wedge x_3$

$x_1 \wedge x_2 \wedge x_4$

$x_1 \wedge x_3 \wedge x_4$

$x_2 \wedge x_3 \wedge x_4$

$x_1 \wedge x_2 \wedge x_3 \wedge x_4$

Hypothesis Space (2)

1	0	0	1	0	0
2	0	1	0	0	0
3	0	0	1	1	1
4	1	0	0	1	1
5	0	1	1	0	0
6	1	1	0	0	0
7	0	1	0	1	0

Simple Rules: There are only 16 simple **conjunctive rules**

of the form $y = x_i \wedge x_j \wedge x_k$

Rule Counterexample

$y=c$

x_1 1100 0

x_2 0100 0

x_3 0110 0

x_4 0101 1

$x_1 \wedge x_2$ 1100 0

$x_1 \wedge x_3$ 0011 1

$x_1 \wedge x_4$ 0011 1

Rule Counterexample

$x_2 \wedge x_3$ 0011 1

$x_2 \wedge x_4$ 0011 1

$x_3 \wedge x_4$ 1001 1

$x_1 \wedge x_2 \wedge x_3$ 0011 1

$x_1 \wedge x_2 \wedge x_4$ 0011 1

$x_1 \wedge x_3 \wedge x_4$ 0011 1

$x_2 \wedge x_3 \wedge x_4$ 0011 1

$x_1 \wedge x_2 \wedge x_3 \wedge x_4$ 0011 1

No simple rule explains the data. The same is true for **simple clauses**.

Hypothesis Space (3)

m-of-n rules: There are 32 possible rules of the form " $y = 1$ if and only if at least m of the following n variables are 1"

1	0	0	1	0	0
2	0	1	0	0	0
3	0	0	1	1	1
4	1	0	0	1	1
5	0	1	1	0	0
6	1	1	0	0	0
7	0	1	0	1	0

Notation: 2 variables from the set on the left. **Value:** Index of the counterexample.

variables 1-of 2-of 3-of 4-of

{X1}

{X2}

{X3}

{X4}

{X1,X2}

{X1, X3}

{X1, X4}

{X2,X3}

variables 1-of 2-of 3-of 4-of

{X2, X4}

{X3, X4}

{X1,X2, X3}

{X1,X2, X4}

{X1,X3,X4}

{X2, X3,X4}

{X1, X2, X3,X4}

Hypothesis Space (3)

m-of-n rules: There are 32 possible rules of the form " $y = 1$ if and only if at least m of the following n variables are 1"

1	0	0	1	0	0
2	0	1	0	0	0
3	0	0	1	1	1
4	1	0	0	1	1
5	0	1	1	0	0
6	1	1	0	0	0
7	0	1	0	1	0

Notation: 2 variables from the set on the left. **Value:** Index of the counterexample.

variables	1-of	2-of	3-of	4-of	variables	1-of	2-of	3-of	4-of
{X1}	3	-	-	-	{X2, X4}	2	3	-	-
{X2}	2	-	-	-	{X3, X4}	4	4	-	-
{X3}	1	-	-	-	{X1,X2, X3}	1	3	3	-
{X4}	7	-	-	-	{X1,X2, X4}	2	3	3	-
{X1,X2}	2	3	-	-	{X1,X3,X4}	1	***	3	-
{X1, X3}	1	3	-	-	{X2, X3,X4}	1	5	3	-
{X1, X4}	6	3	-	-	{X1, X2, X3,X4}	1	5	3	3
{X2,X3}	2	3	-	-					

Views of Learning

- ❖ Learning is the removal of our remaining uncertainty:
- ❖ Learning requires guessing a good hypothesis class
 - ❖ Start with a small class and enlarge it until it contains an hypothesis that fits the data.
- ❖ We could be wrong !
 - ❖ Our guess of the hypothesis space could be wrong
 - ❖ $y=x^4 \wedge$ one-of (x_1, x_3) is also consistent

General strategies for Machine Learning

- ❖ Develop flexible hypothesis spaces:
 - ❖ Decision trees, neural networks, nested collections.
- ❖ Develop representation languages for restricted classes of functions:
 - ❖ E.g., Functional representation (n-of-m); Grammars; linear functions; stochastic models

General strategies for Machine Learning

- ❖ Develop flexible hypothesis spaces:
 - ❖ Decision trees, neural networks, nested collections.
- ❖ Develop representation languages for restricted classes of functions:

In either case:


- ❖ Develop algorithms for finding a hypothesis in our hypothesis space, that fits the data
- ❖ And hope that they will generalize well

Hypothesis Space -- Real-Value Features


WHAAAA?!?!?




Example problem




Hypothesis space:



Hypothesis space: linear model



Hypothesis space: linear model



w, b are the parameters to represent a linear function